

# Predict a distribution model

D'Souza, Ajay  
ajaydsouza@gatech.edu

## **Abstract**

Fit a model to predict the mean and variance of a distribution given two predictor variables. We are provided with a training set with a sample of 200 values for every set of predictor variables from which the model needs to learn to predict the  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$

# Contents

List of Figures	4
List of Tables	5
1 Introduction	6
2 Problem Definition	6
3 Exploratory Data Analysis	6
3.1 Distribution and Density Plots . . . . .	6
3.2 Correlation Analysis . . . . .	8
3.3 Mean(Y)- $\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ Plots with $X_1, X_2$ . . . . .	10
3.4 $\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ Quartile plots with $X_1, X_2$ . . . . .	13
4 Proposed Method	14
4.1 Predicting Mean(Y)- $\hat{\mu}$ . . . . .	14
4.1.1 Logistic Regression . . . . .	14
4.1.2 Logistic Regression with Polynomial terms . . . . .	14
4.1.3 Linear Models . . . . .	14
4.2 Predicting Variance of Y - $\hat{\sigma}$ . . . . .	15
4.2.1 Logistic Regression . . . . .	16
4.2.2 Logistic Regression with Polynomial terms . . . . .	16
4.3 Linear Models . . . . .	16
4.4 10 Fold Cross Validation for choosing model parameters . . . . .	17
4.5 Evaluating models with Bootstrapping with B=100 Iterations . . . . .	17
5 Results and Observations	18
5.1 Predicting Mean of Y - $\hat{\mu}$ . . . . .	18
5.1.1 Logistic Regression . . . . .	18
5.1.2 Logistic Regression with Polynomial Terms . . . . .	21
5.1.3 Linear Methods . . . . .	25
5.2 Predicting Variance of Y - $\hat{\sigma}$ . . . . .	33
5.2.1 Logistic Regression . . . . .	33
5.2.2 Logistic Regression with Polynomial Terms . . . . .	35
5.2.3 Linear Methods . . . . .	39

<b>6</b>	<b>Test Results</b>	<b>46</b>
6.1	Cross Validation Results . . . . .	46
6.2	Bootstrapping 100 Iterations Test Results . . . . .	46
6.3	T-test and Wilcox Test for $\text{Mean}(Y) - \hat{\mu}$ prediction models . . . . .	50
6.4	T-test and Wilcox Test for $\text{Variance}(Y) - \hat{\sigma}$ prediction models . . . . .	50
<b>7</b>	<b>Conclusions</b>	<b>51</b>
7.1	$\text{Mean}(Y) - \hat{\mu}$ . . . . .	51
7.2	$\text{Variance}(Y) - \hat{\sigma}$ . . . . .	52
<b>8</b>	<b>Appendix</b>	<b>52</b>
<b>9</b>	<b>References</b>	<b>95</b>

## List of Figures

1	Box Plot and Density Plot of Training Data $X_1, X_2, \text{Mean}(Y)-\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ . . . . .	7
2	Density Plot of Training Data Y for Specific $X_1, X_2$ . . . . .	8
3	Correlation Plot $X_1, X_2, \text{Mean}(Y)-\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ with correlation values . . . . .	9
4	Detailed Correlation $X_1, X_2, \text{Mean}(Y)-\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ . . . . .	11
5	Mean(Y)- $\hat{\mu}$ Versus Variance(Y)- $\hat{\sigma}$ Plot . . . . .	12
6	$X_1, X_2$ , Versus Mean(Y)- $\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ Quartile Plots . . . . .	13
7	Logistic Regression Residuals for Mean(Y)- $\hat{\mu}$ colored by Mean(Y)- $\hat{\mu}$ Quartile	20
8	Cross Validation MSE for predicting Mean(Y)- $\hat{\mu}$ plotted against Predictor Combinations . . . . .	21
9	Logistic Regression Residuals for Mean(Y)- $\hat{\mu}$ colored by Mean(Y)- $\hat{\mu}$ Quartile	24
10	Linear Regression Residuals for Mean(Y)- $\hat{\mu}$ colored by Mean(Y)- $\hat{\mu}$ Quartile	26
11	Linear Regression (Polynomial) Residuals for Mean(Y)- $\hat{\mu}$ colored by Mean(Y)- $\hat{\mu}$ Quartile . . . . .	29
12	LASSO (Polynomial) Cross Validation Plot for Mean(Y)- $\hat{\mu}$ for choosing $\lambda$	30
13	LASSO (Polynomial) Residuals for Mean(Y)- $\hat{\mu}$ colored by Mean(Y)- $\hat{\mu}$ Quartile . . . . .	32
14	Logistic Regression Residuals for Variance(Y)- $\hat{\sigma}$ colored by Variance(Y)- $\hat{\sigma}$ Quartile . . . . .	34
15	Cross Validation MSE for predicting Variance(Y)- $\hat{\sigma}$ plotted against Predictor Combinations . . . . .	35
16	Logistic Regression Residuals for Variance(Y)- $\hat{\sigma}$ colored by Variance(Y)- $\hat{\sigma}$ Quartile . . . . .	38
17	Linear Regression Residuals for Variance(Y)- $\hat{\sigma}$ colored by Variance(Y)- $\hat{\sigma}$ Quartile . . . . .	40
18	Linear Regression (Polynomial) Residuals for Variance(Y)- $\hat{\sigma}$ colored by Variance(Y)- $\hat{\sigma}$ Quartile . . . . .	42
19	LASSO (Polynomial) Cross Validation Plot for Variance(Y)- $\hat{\sigma}$ for choosing $\lambda$	43
20	LASSO (Polynomial) Residuals for Variance(Y)- $\hat{\sigma}$ colored by Variance(Y)- $\hat{\sigma}$ Quartile . . . . .	45
21	Box-plot of Bootstrapping MSE results of both Mean(Y)- $\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ predictor models . . . . .	47
22	Bootstrapping 100 Iterations - Average MSE , Variance MSE for the Mean(Y)- $\hat{\mu}$ prediction models . . . . .	48

23	Bootstrapping 100 Iterations - Average MSE , Variance MSE for the Variance(Y)- $\hat{\sigma}$ prediction models . . . . .	49
----	--	----

**List of Tables**

1	. . . . .	6
2	. . . . .	46
3	. . . . .	46
4	. . . . .	46
5	. . . . .	50
6	. . . . .	51

# 1 Introduction

Given two predictor variables  $X_1$  and  $X_2$ , they emit a response  $Y$ . The distribution of  $Y$  given  $X_1, X_2$  is not provided. We are given a set of values of  $Y$  for a range of specific combinations of  $X_1, X_2$ . We need to learn from this data and fit a appropriate model to predict both the  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$  of the distribution given a specific value of  $X_1, X_2$ . The goal is to learn from this data and develop a model that will estimate this mean  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$  for any given  $X_1, X_2$

# 2 Problem Definition

We are provided with a data set that is  $2911 \times 202$ . The structure of the data is as shown in table (1). Each row consists of  $X_1, X_2, Y_1 \dots Y_{200}$  responses. The distribution of response  $Y$  given a  $X_1, X_2$  is not provided and is not assumed. From this data we compute the mean and variance of  $Y_1 \rightarrow Y_{200}$  for each  $X_1, X_2$  as  $\text{Mean}(Y) - \hat{\mu} = \text{mean}(Y_1 \rightarrow Y_{200})$  and variance as  $\text{Variance}(Y) - \hat{\sigma} = \text{Variance}(Y_1 \rightarrow Y_{200})$ .

$X_1$	$X_2$	$Y_1 \rightarrow Y_{200}$	$\text{muhat} = \text{Mean}(Y) - \hat{\mu}$	$\text{Vhat} = \text{Variance}(Y) - \hat{\sigma}$
1	2	2 $\rightarrow$ 202	computed	computed

Table 1: Structure of the Data

# 3 Exploratory Data Analysis

## 3.1 Distribution and Density Plots

- Figure (1) shows the distribution of the training data along with the density plots for the  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$ . The density plots for  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$  are for all values of  $X_1, X_2$ . Since are more interested in the distribution of  $Y$  given a specific  $X_1, X_2$ . Figure (2) are some density plots of  $Y$  for different specific values of  $X_1, X_2$ .

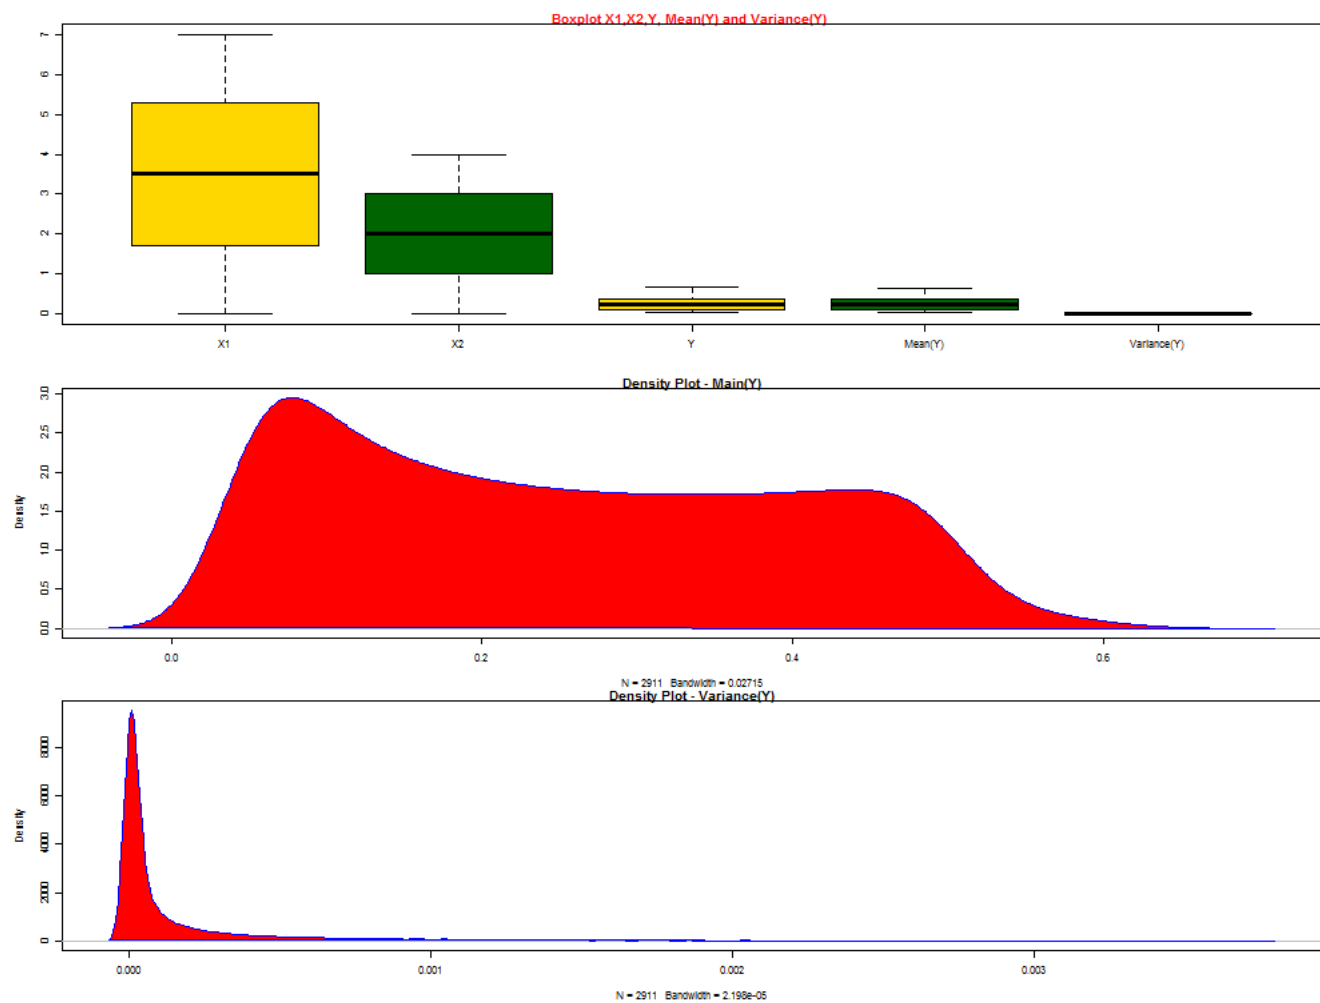


Figure 1: Box Plot and Density Plot of Training Data  $X_1, X_2, \text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$

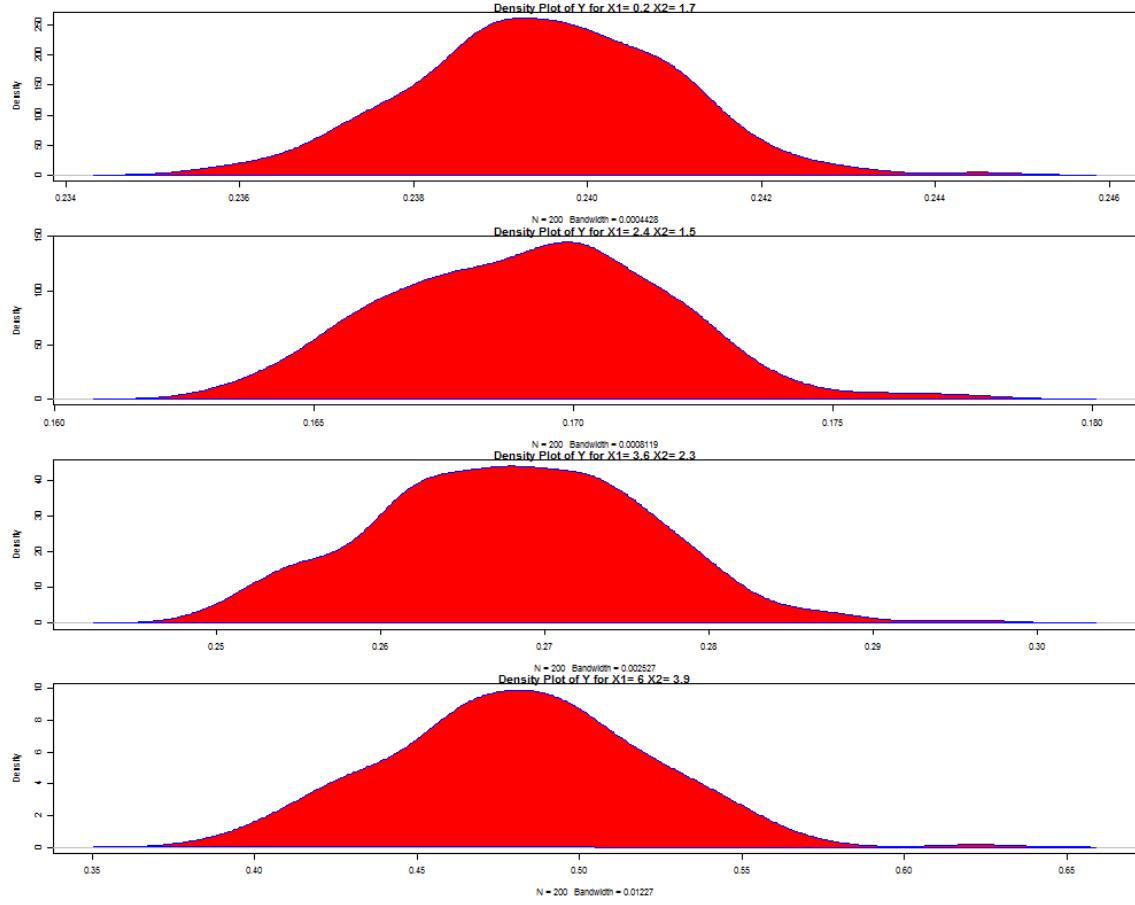


Figure 2: Density Plot of Training Data Y for Specific  $X_1, X_2$

### 3.2 Correlation Analysis

- From figure (3)  $\text{Mean}(Y) - \hat{\mu}$  is strongly correlated to  $X_2$  with a correlation of  $\approx 0.99$ .  $\text{Variance}(Y) - \hat{\sigma}$  has a moderate correlation to both  $X_1, X_2$  of  $\approx 0.5$ .



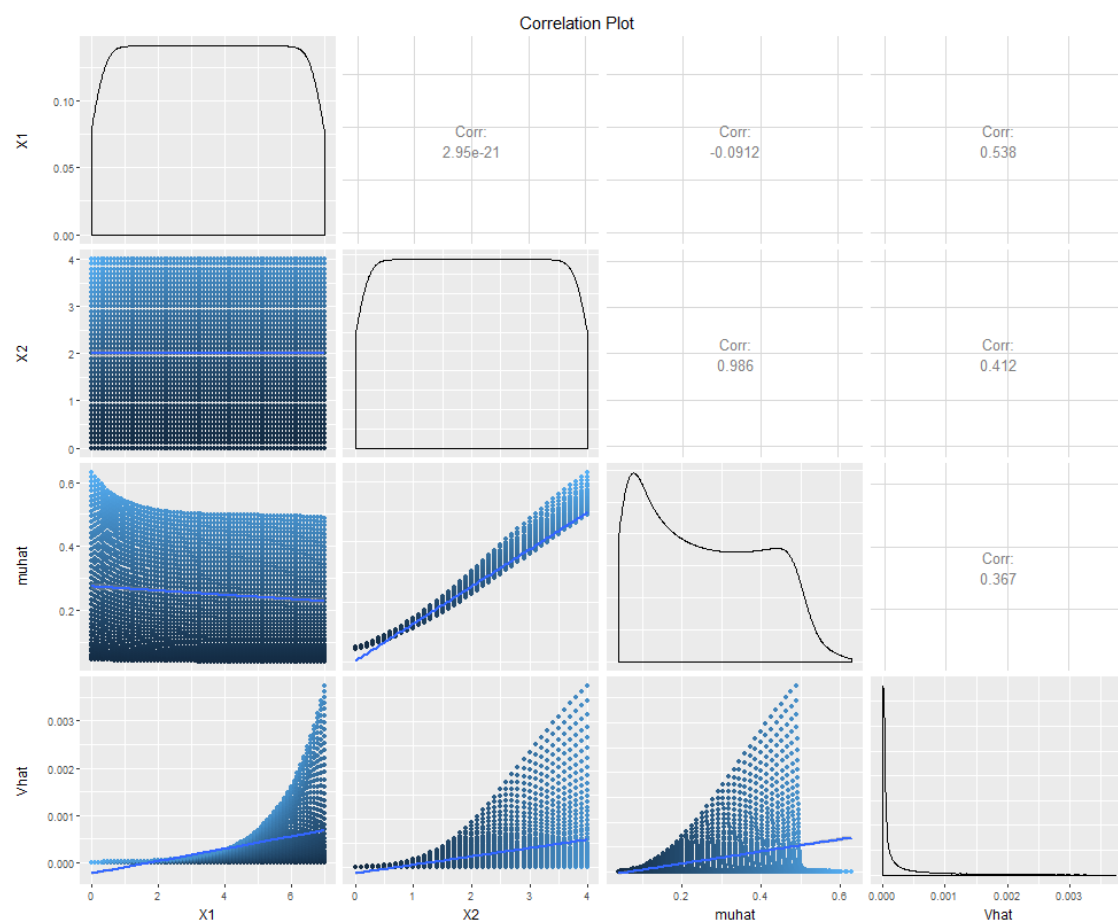


Figure 3: Correlation Plot  $X_1, X_2, \text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$  with correlation values

### 3.3 Mean(Y)- $\hat{\mu}$ and Variance(Y)- $\hat{\sigma}$ Plots with $X_1, X_2$

- Figure (4) plots the Mean(Y)- $\hat{\mu}$  and Variance(Y)- $\hat{\sigma}$  for  $X_1$  for each value of  $X_2$  and similarly for  $X_2$  for each value of  $X_1$ . Mean(Y)- $\hat{\mu}$  appears to be independent of  $X_1$  except at low values of  $X_1$  between 0 . . . 2. This correlation with  $X_1$  is negative and very low. But correlation of Mean(Y)- $\hat{\mu}$  with  $X_2$  is very strong and increases almost linearly with  $X_2$ . Variance(Y)- $\hat{\sigma}$  increases rapidly beyond  $X_1 > 5$  and  $X_2 > 2$ . For  $X_1, X_2$  values in a range lower than this Variance(Y)- $\hat{\sigma}$  remains low. The plot of Mean(Y)- $\hat{\mu}$  versus  $X_2$  and the plot of Variance(Y)- $\hat{\sigma}$  versus  $X_1$  and  $X_2$  shows patterns of a sigmoid curve.

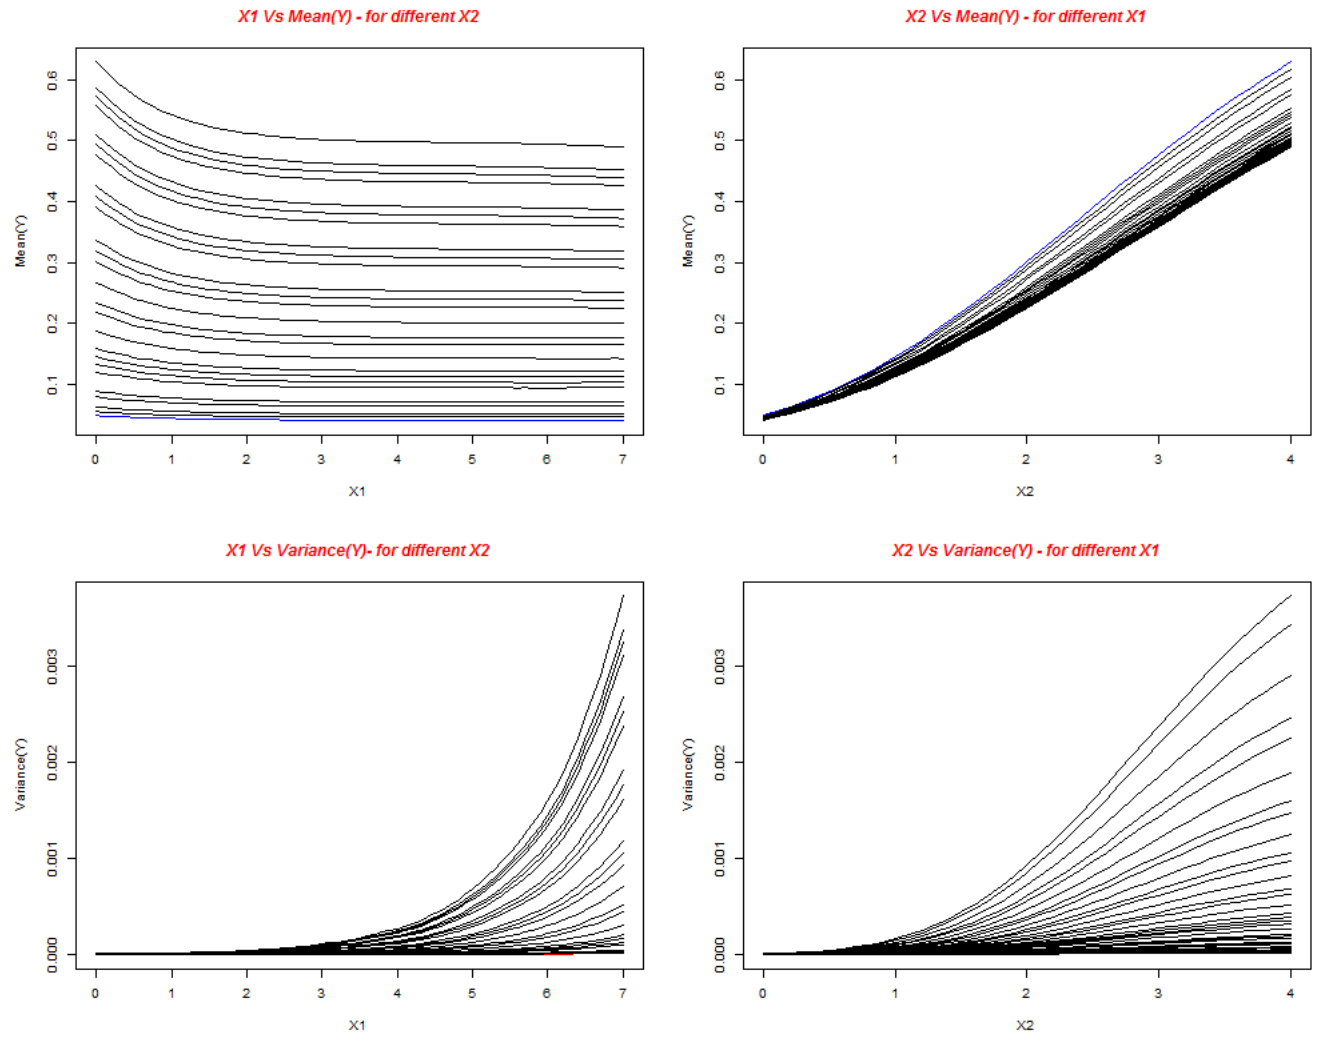


Figure 4: Detailed Correlation  $X_1, X_2, \text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$

- Figure (5) is an interesting plot that gives us an hint on the methods that could be used for predicting  $\text{Mean}(Y) - \hat{\mu}$ . It plots  $\text{Mean}(Y) - \hat{\mu}$  versus the variance  $\text{Variance}(Y) - \hat{\sigma}$ . The variance is not constant but increases and shows a wide spread as  $\text{Mean}(Y) - \hat{\mu}$  increases. This indicates heteroscedasticity.

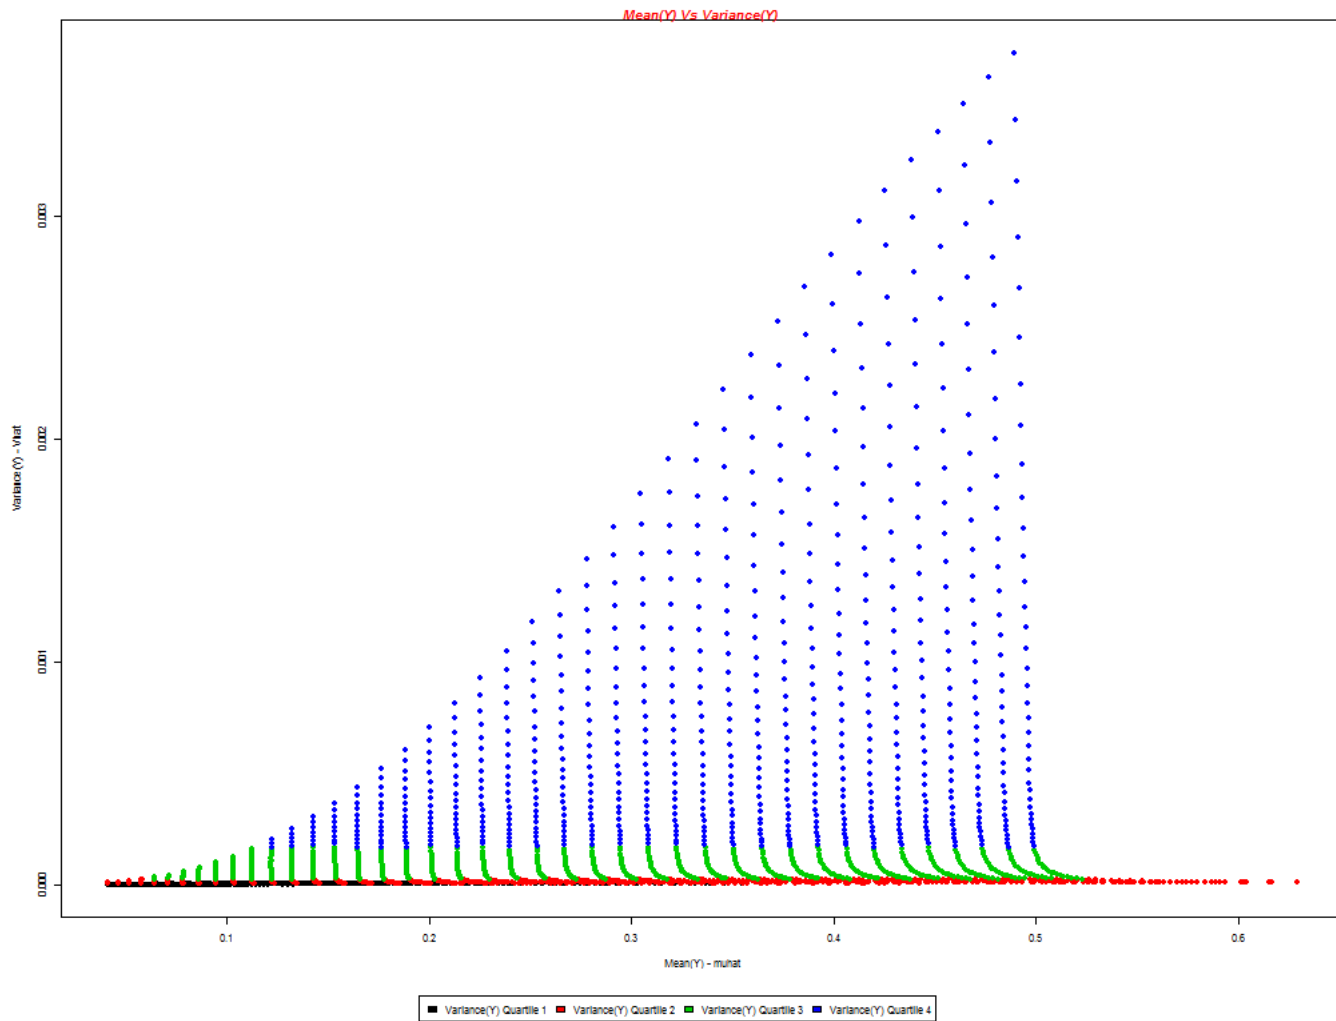


Figure 5:  $\text{Mean}(Y) - \hat{\mu}$  Versus  $\text{Variance}(Y) - \hat{\sigma}$  Plot

### 3.4 $\hat{\mu}$ and $\text{Variance}(Y)-\hat{\sigma}$ Quartile plots with $X_1, X_2$

- Figure (6) gives a better picture of the distribution as it plots the Quartiles of  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$  simultaneously for  $X_1, X_2$ . The  $\text{Mean}(Y)-\hat{\mu}$  quartile bands change very little with  $X_1$ , but increase almost linearly with  $X_2$ . But the  $\text{Variance}(Y)-\hat{\sigma}$  quartile bands increase for both  $X_1, X_2$ .

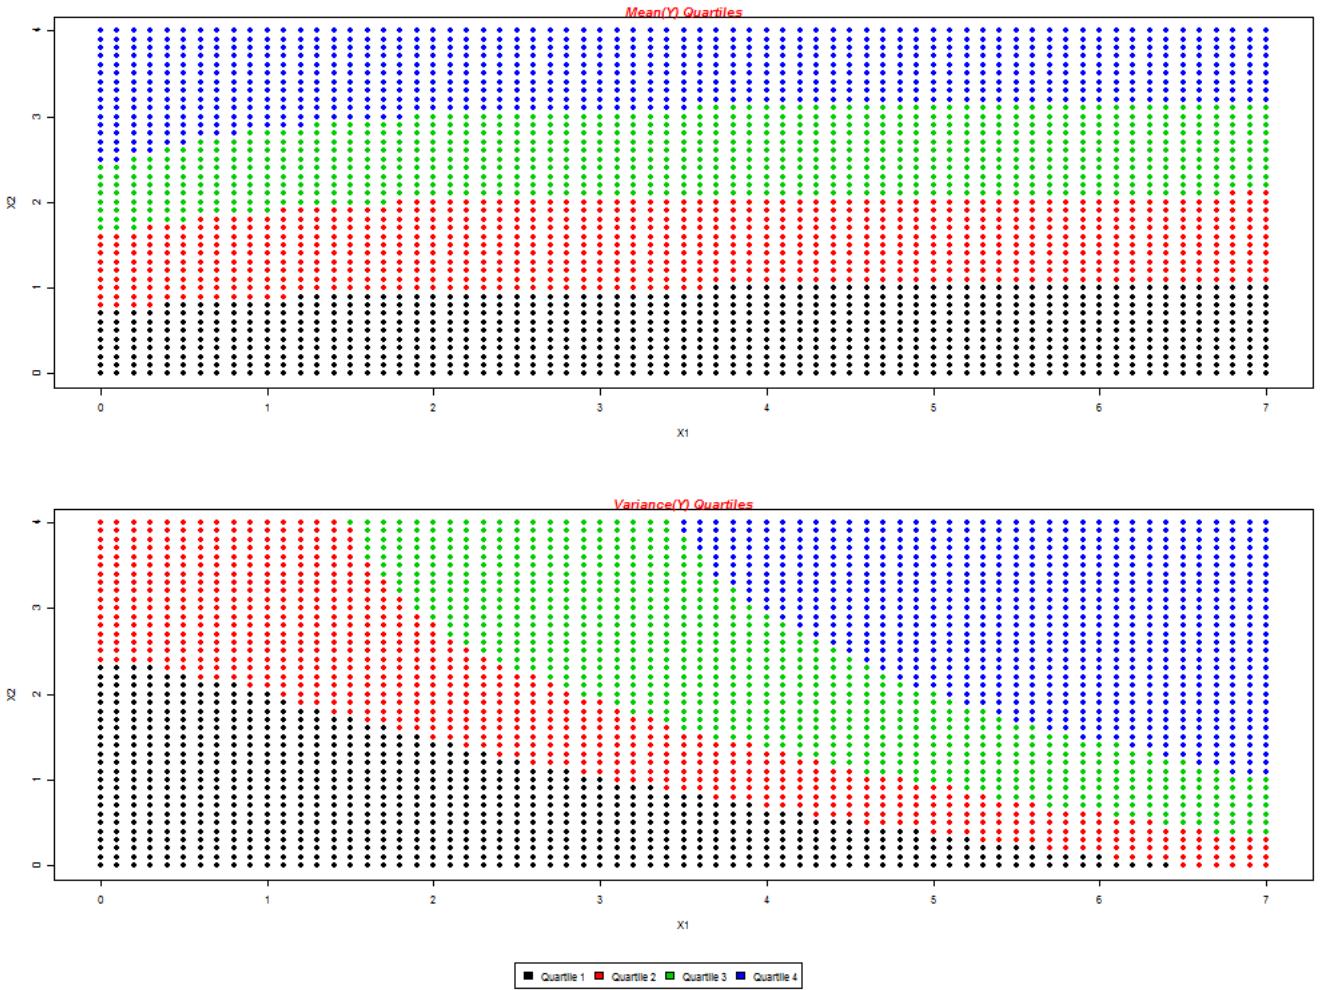


Figure 6:  $X_1, X_2$ , Versus  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$  Quartile Plots

## 4 Proposed Method

### 4.1 Predicting Mean(Y)- $\hat{\mu}$

From plot (5) we see that variance is heteroscedastic and increases and widens in spread as Mean(Y)- $\hat{\mu}$  increases. Ordinary Least Squares based linear regression methods are not the best models to use in such a case as they assume homoscedasticity. We also know that Mean(Y)- $\hat{\mu}$  is in the range of 0...1. The value of Mean(Y)- $\hat{\mu}$  can be taken as P(x), the probability of success. Logistic Regression has the advantage that it does not make assumptions on the uniformity of variance nor the nature of distribution. From the plot in (4) of Mean(Y)- $\hat{\mu}$  versus  $X_2$  to which Mean(Y)- $\hat{\mu}$  is highly correlated, we can see it follows a pattern that fits the sigmoid function. So Logistic Regression model which models the sigmoid function for the response makes a nice fit for prediction of Mean(Y)- $\hat{\mu}$  based on  $X_1, X_2$ . Thus based on the pattern of data, Logistic Regression appears to be the best model for predicting Mean(Y)- $\hat{\mu}$ . We will also explore the Linear Models so we can compare the results and choose the best model.

#### 4.1.1 Logistic Regression

Here we fit a standard Logistic Regression model with Mean(Y)- $\hat{\mu}$  as response and  $X_1, X_2, X_1 * X_2$  as the predictor variables.

#### 4.1.2 Logistic Regression with Polynomial terms

We explore the Logistic Regression model further by performing a ten fold cross validation to choose the best combination of interaction and polynomial terms of  $X_1, X_2$ . We will explore up-to the 6<sup>th</sup> degree of polynomial for each predictor variable  $X_1 * X_2$ . For each combination we will get the Cross Validation Average MSE and choose the polynomial Logistic model with the lowest Cross Validation MSE. The selected model will then be trained on the whole training data.

#### 4.1.3 Linear Models

From Figure (6), Mean(Y)- $\hat{\mu}$  is highly positively correlated to  $X_2$  and very slightly negatively correlated to low values of  $X_1$ . Though Logistic Regression is the first choice for predicting Mean(Y)- $\hat{\mu}$  due to reasons explained in the previous section, we will still evaluate Linear Regression models if it can add value due to this strong correlation. We will evaluate the performance of with a regular linear regression model, a linear regression model with polynomial terms and a LASSO model with polynomial terms. The model

parameters will be selected using 10 fold cross validation on the training data. Model performance will be compared using bootstrapping to select the best model for generating the test results

**Linear Regression** This is base model, where we try to fit a linear regression model for predicting  $\text{Mean}(Y) - \hat{\mu}$  using predictor variables  $X_1, X_2, X_1 * X_2$

**Linear Regression with Polynomial Terms** Here we will use ten fold cross validation to choose the best polynomial degree for each of  $X_1$  and  $X_2$  for the linear regression model for predicting  $\text{Mean}(Y) - \hat{\mu}$ . We propose to perform cross validation for every Polynomial combination of  $X_1$  and  $X_2$  to a maximum degree of 6. That gives us 36 models to choose from using ten fold cross validation for each. The model with the lowest MSE on the cross validation set is chosen. The chosen model is trained on the whole training data.

**LASSO with Polynomial Terms** Having linear regression with higher order polynomial terms might work well during training, but it could be picking up noise and over fitting the training data. One way to get around this while still trying to take advantage of polynomial terms to learn non linear relationships is to use a regularization model like the LASSO. LASSO has the optimization goal of minimizing the coefficients built into the learning process for building the model and so will try to shrink terms that contribute less to the model.

Here we use ten fold cross validation to again choose the best polynomial terms of  $X_1$  and  $X_2$  with best  $\lambda$  for LASSO. The best cross validated LASSO linear regression model with the chosen polynomial terms in  $X_1$  and  $X_2$  is then trained on the whole training data with a further cross validation done to choose the best  $\lambda$  for LASSO. The model with the chosen polynomial terms and  $\lambda$  is then trained on the whole training set to get the optimum LASSO Linear Regression model.

## 4.2 Predicting Variance of $Y - \hat{\sigma}$

From the plots (5) we see heteroscedasticity since  $\text{Variance}(Y) - \hat{\sigma}$  increases and its distribution spreads with  $\text{Mean}(Y) - \hat{\mu}$  which in turn depends on  $X_1, X_2$ . The OLS based linear regression models assume homoscedasticity and are not a natural fit to model this pattern. From the density plots of the  $\text{Variance}(Y) - \hat{\sigma}$  in (1) we see that  $\text{Variance}(Y) - \hat{\sigma}$  has a skewed distribution with the peak very close to 0 with a rapid fall in peak for values  $> 0$ . We know that variance cannot be negative and in this case we know  $\text{Variance}(Y) - \hat{\sigma}$  is  $0 \dots 1$ . So we have to fit a model where this constraint can be enforced. Any regular

linear regression model fitted without enforcing this constraint might give negative values for  $\text{Variance}(\hat{Y}) - \hat{\sigma}$ . Negative predictions are meaningless. The plots of the  $\text{Variance}(\hat{Y}) - \hat{\sigma}$  versus  $X_1$  and  $X_2$  in (4) shows a pattern which is similar to the sigmoid function. Logistic Regression Model seeks to model the sigmoid function. Thus for predicting  $\text{Variance}(\hat{Y}) - \hat{\sigma}$  Logistic Regression Model seems to be good first option. We will also explore the Linear Models so we can compare their performance with the Logistic Regression Models to choose the best model to be used for the test prediction of  $\text{Variance}(\hat{Y}) - \hat{\sigma}$ . We will seek to fit the following models for prediction of  $\text{Variance}(\hat{Y}) - \hat{\sigma}$

#### 4.2.1 Logistic Regression

We take  $\text{Variance}(\hat{Y}) - \hat{\sigma}$  as the  $P(X)$  value and fit a standard logistic regression model with  $X_1, X_2, X_1 * X_2$  as the predictors

#### 4.2.2 Logistic Regression with Polynomial terms

We explore the Logistic Regression model further by performing a ten fold cross validation to choose the best combination of interaction and polynomial terms of  $X_1, X_2$ . We will explore up-to the 6<sup>th</sup> degree of polynomial for each predictor variable  $X_1 * X_2$ . For each combination we will get the Cross Validation Average MSE and from that choose the model with the lowest Cross Validation MSE.

### 4.3 Linear Models

From Figure (6),  $\text{Variance}(\hat{Y}) - \hat{\sigma}$  is correlated to both  $X_1$  and  $X_2$ . Though Logistic Regression is the first choice for predicting  $\text{Variance}(\hat{Y}) - \hat{\sigma}$  due to reasons explained in the previous section, we will still evaluate the following Linear Regression models. The model parameters will be selected using 10 fold cross validation on the training data. Model performance will be compared using bootstrapping to select the best model for generating the test results

**Linear Regression** We take  $\text{Variance}(\hat{Y}) - \hat{\sigma}$  as the response and perform a standard linear regression with  $X_1, X_2, X_1 * X_2$  as the predictors

**Linear Regression with Polynomial Terms** We will use ten fold cross validation to choose the best polynomial degree for each of  $X_1$  and  $X_2$  for the linear regression model for predicting  $\text{Variance}(\hat{Y}) - \hat{\sigma}$ . We propose to perform cross validation for every Polynomial combination of  $X_1$  and  $X_2$  to a maximum degree of 6. That gives us 36 models to



choose from using ten fold cross validation for each. The model with the lowest MSE on the cross validation set is chosen. The chosen model is trained on the whole training data.

**LASSO with Polynomial Terms** Using  $\text{Variance}(\hat{\sigma})$  as the response variable we will train the LASSO model using ten fold cross validation. The first cross validation is to choose the best polynomial terms for  $X_1$  and  $X_2$  for the LASSO model. Once the optimum polynomial terms are chosen we perform another round of cross validation to get the best  $\lambda$ . The chosen model is then trained on the whole training data.

#### 4.4 10 Fold Cross Validation for choosing model parameters

1. The models are trained and optimum parameters for the model are chosen using the process of 10 fold cross validation.
2. For this purpose the training data is divided into random set of 10 folds
3. For cross validation all combinations of tunable parameters will be used to train the model using each set of 9 folds as the training set and the remaining 1 fold as the test set.
4. The average MSE across each of the 10 test folds is taken for every combination of parameters
5. The model with the combination of parameters with the lowest average error rate is chosen as the optimum one for a given method
6. The chosen model is trained using these optimum parameters on the whole training set
7. The performance of the model is evaluated using bootstrapping as described next

#### 4.5 Evaluating models with Bootstrapping with B=100 Iterations

Once we have generated the different models described here, we evaluate the model using the sampling method of bootstrapping. We use all the training data for bootstrapping and perform iterations for 100 samples with replacement as follows.

- A robust way to test the models is using bootstrapping. We can perform this with  $B = 100$  iterations
- Since we have sufficient data of 2911 rows, we can pick a random 40% of the data with replacement as the held back test data during each bootstrapping cycle.
- Training the model on the balance 60% of the data for each cycle will reduce the correlation between the models trained and thus reduce the variance on the MSE from the bootstrapping process.
- Using the optimum parameters we have already chosen we train each model for Mean(Y)- $\hat{\mu}$  and Variance(Y)- $\hat{\sigma}$  on the 60% data designated as training data in each bootstrapping cycle. For each of the trained models we then compute the MSE as described in (1) and (2) on the 40% data chosen as the test data. The results are tabulated for each cycle

$$MSE_{\hat{\mu}}^{model} = \frac{1}{n} \sum_{i=1}^{i=n} \left( \hat{\mu}_i^{test} - \hat{f}_{\hat{\mu}}^{model} (x_i^{test}) \right)^2 \quad (1)$$

$$MSE_{\hat{\sigma}}^{model} = \frac{1}{n} \sum_{i=1}^{i=n} \left( \hat{\sigma}_i^{test} - \hat{f}_{\hat{\sigma}}^{model} (x_i^{test}) \right)^2 \quad (2)$$

- Once the 100 bootstrapping iterations are completed, the mean and variance of these  $MSE_{\hat{\mu}}^{model}$  and  $MSE_{\hat{\sigma}}^{model}$  results from each cycle for each model is computed
- We can now perform a statistical test like T Test or a W Test to reliably choose the best model
- The best model is used to predict the results on the test data set provided

## 5 Results and Observations

The following are the results from cross validation and training the models.

### 5.1 Predicting Mean of Y - $\hat{\mu}$

#### 5.1.1 Logistic Regression

- The very low P-value for  $X_2$  shows that  $X_2$  is important for predicting the response Mean(Y)- $\hat{\mu}$  in this model. As expected the  $X_2$  has a positive coefficient that is

approximately 20 times larger in magnitude than the negative coefficient of  $X_1$ . So  $X_2$  has a larger contribution to the response  $\text{Mean}(Y) - \hat{\mu}$ .

```
> summary(lr.mu.fit)
```

Call:

```
glm(formula = muhat ~ X1 * X2, family = binomial(logit), data = data0.train)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-0.116587	-0.039413	0.001049	0.027583	0.151184

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-2.670132	0.228276	-11.697	<2e-16 ***
X1	-0.021026	0.057231	-0.367	0.713
X2	0.763357	0.085642	8.913	<2e-16 ***
X1:X2	-0.007681	0.021305	-0.361	0.718

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 343.2493 on 2764 degrees of freedom  
 Residual deviance: 6.1953 on 2761 degrees of freedom  
 AIC: 1690.1

Number of Fisher Scoring iterations: 5

- Figure (7) plots the fitted values for  $\text{Mean}(Y) - \hat{\mu}$  against the true values. The figure also plots the standardized residuals from the model with  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$ . The plots are colored to indicate the quartiles to which the predicted  $\text{Mean}(Y) - \hat{\mu}$  values belong. Ideally we should have a  $45^\circ$  line for the true and fitted values, in this case plot appears to have a spread as the value of  $\text{Mean}(Y) - \hat{\mu}$  increases. The standardized residuals also spread higher for upper quartiles of  $\text{Mean}(Y) - \hat{\mu}$ .

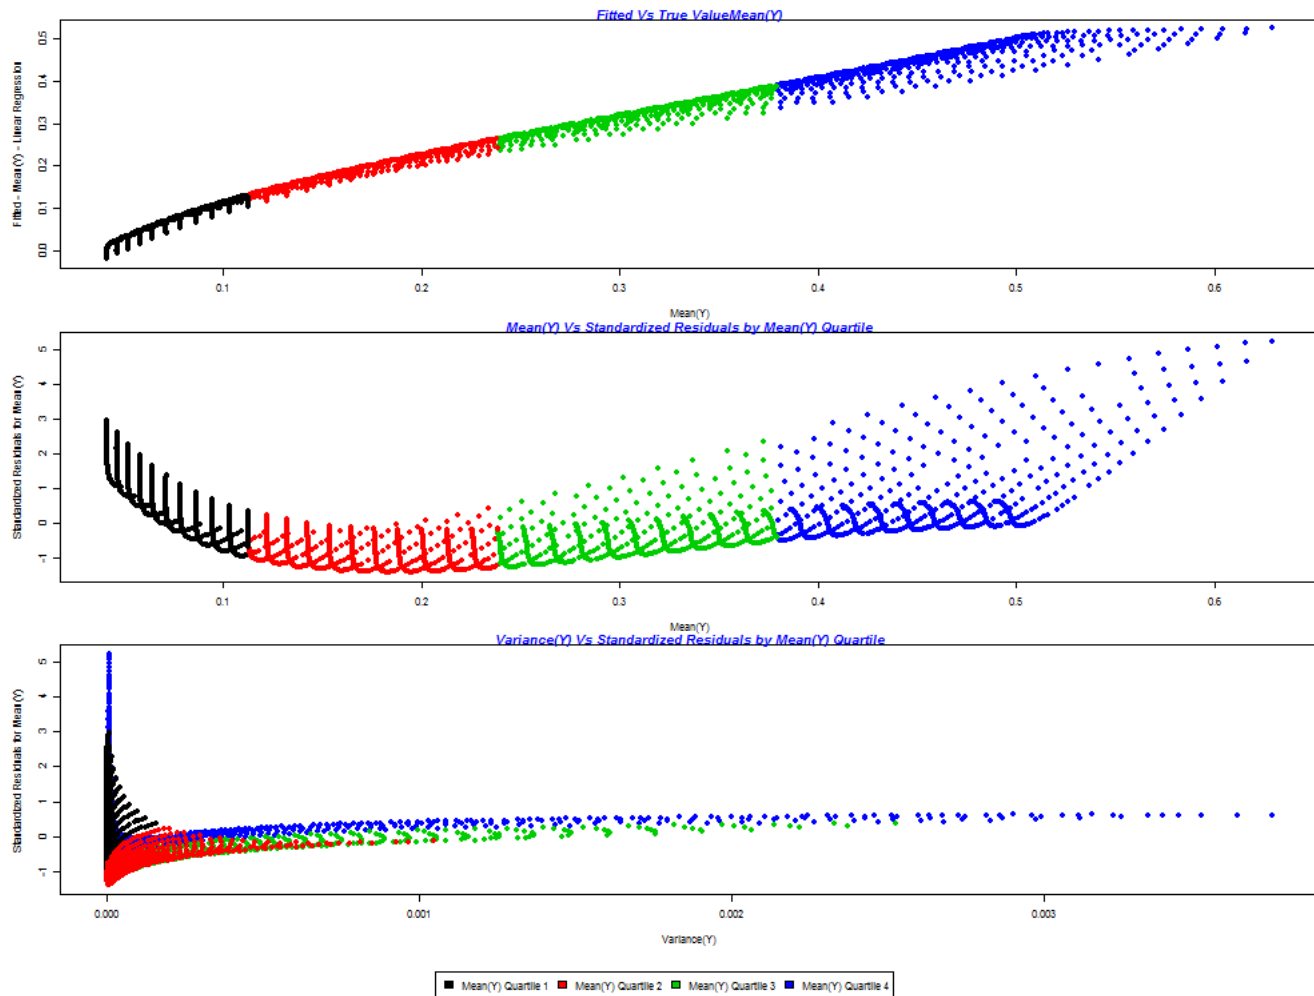


Figure 7: Logistic Regression Residuals for  $\text{Mean}(Y) - \hat{\mu}$  colored by  $\text{Mean}(Y) - \hat{\mu}$  Quartile

### 5.1.2 Logistic Regression with Polynomial Terms

- The 10 fold cross validation plot for MSE for various combinations of polynomial degree terms of  $X_1, X_2$  is shown in figure (8). The ten fold cross validation picks degree 6 for the interaction of  $X_1, X_2$  as the best model.

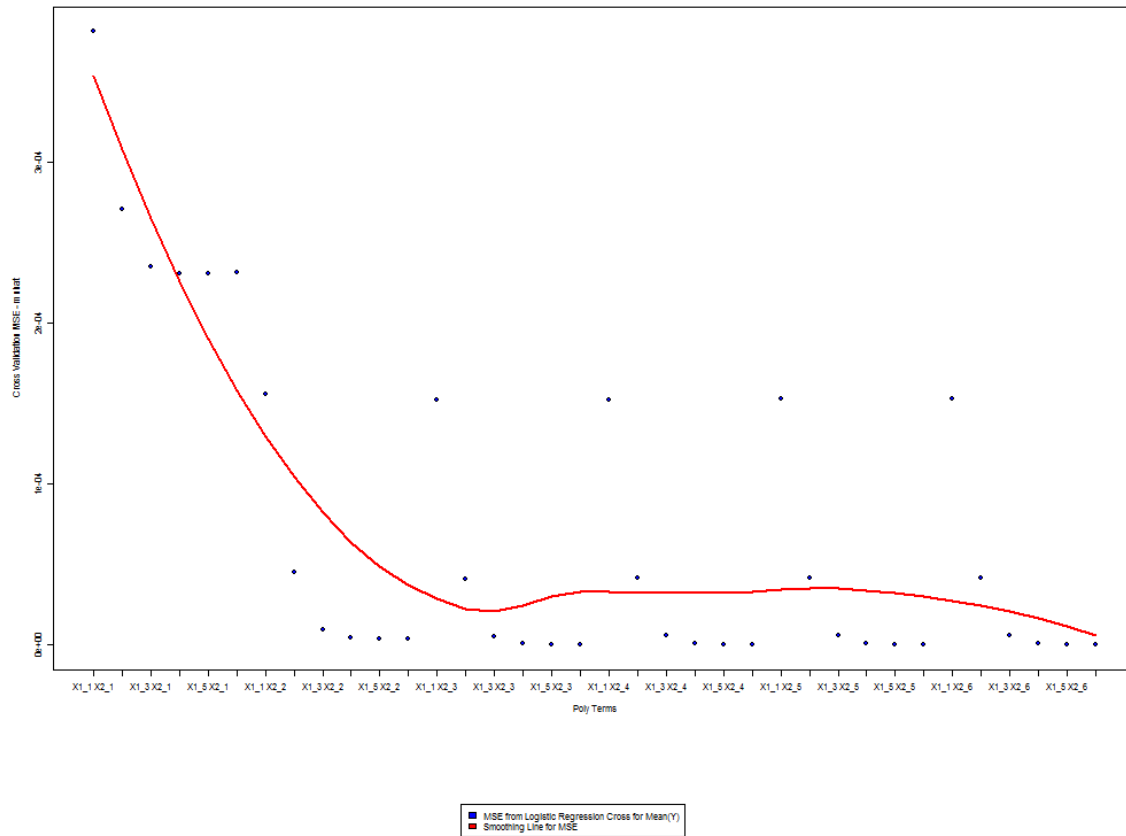


Figure 8: Cross Validation MSE for predicting  $\text{Mean}(Y) - \hat{\mu}$  plotted against Predictor Combinations

- For the best model chosen, the very low P-value for  $X_2$  shows that  $X_2$  is important for predicting the response  $\text{Mean}(Y) - \hat{\mu}$  in this model. As expected the  $X_2$  has the largest positive coefficient, so  $X_2$  has a larger contribution to the response  $\hat{\mu}$ . The model also has a very low residual deviance of  $1.1517e-04$ . This indicates a good fit.

```

> summary(lrp.mu.fit)

Call:
glm(formula = muhat ~ poly(X1, poly.lrp.muhat.min.x1) * poly(X2,
  poly.lrp.muhat.min.x2), family = binomial(logit), data = data0.train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.010e-03  -9.421e-05  2.210e-06   9.125e-05   8.068e-04

Coefficients:
                                Estimate Std. Error
(Intercept)                    -1.305596    0.095596
poly(X1, poly.lrp.muhat.min.x1)1 -3.796362    2.821445
poly(X1, poly.lrp.muhat.min.x1)2  2.571445    2.821445
poly(X1, poly.lrp.muhat.min.x1)3 -1.365234    2.831445
poly(X1, poly.lrp.muhat.min.x1)4  0.496035    2.831445
poly(X1, poly.lrp.muhat.min.x1)5 -0.154637    2.831445
poly(X1, poly.lrp.muhat.min.x1)6  0.043054    2.831445
poly(X2, poly.lrp.muhat.min.x2)1 48.534772    3.221445
poly(X2, poly.lrp.muhat.min.x2)2 -5.942285    3.161445
poly(X2, poly.lrp.muhat.min.x2)3  0.725413    3.131445
poly(X2, poly.lrp.muhat.min.x2)4 -0.048707    3.101445
poly(X2, poly.lrp.muhat.min.x2)5  0.014232    3.011445
poly(X2, poly.lrp.muhat.min.x2)6 -0.005096    2.741445
poly(X1, poly.lrp.muhat.min.x1)1:poly(X2, poly.lrp.muhat.min.x2)1 -52.359326 168.101445
poly(X1, poly.lrp.muhat.min.x1)2:poly(X2, poly.lrp.muhat.min.x2)1 32.558110 167.781445
.
.
poly(X1, poly.lrp.muhat.min.x1)2:poly(X2, poly.lrp.muhat.min.x2)6  0.029088 142.901445
poly(X1, poly.lrp.muhat.min.x1)3:poly(X2, poly.lrp.muhat.min.x2)6 -0.012942 143.541445
poly(X1, poly.lrp.muhat.min.x1)4:poly(X2, poly.lrp.muhat.min.x2)6  0.009426 143.311445
poly(X1, poly.lrp.muhat.min.x1)5:poly(X2, poly.lrp.muhat.min.x2)6 -0.012221 143.211445
poly(X1, poly.lrp.muhat.min.x1)6:poly(X2, poly.lrp.muhat.min.x2)6 -0.003717 143.421445
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 3.4325e+02  on 2764  degrees of freedom
Residual deviance: 1.1517e-04  on 2716  degrees of freedom

```

AIC: 1781.5

Number of Fisher Scoring iterations: 6

- Figure (9) plots the fitted values for  $\text{Mean}(Y)-\hat{\mu}$  against the true values. The figure also plots the standardized residuals from the model with  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$ . The plots are colored to indicate the quartiles to which the predicted  $\text{Mean}(Y)-\hat{\mu}$  values belong. Ideally we should have a  $45^\circ$  line for the true and fitted values, in this case the line appears to be almost ideal. The standardized residuals show a slight spread higher for upper quartiles of  $\text{Mean}(Y)-\hat{\mu}$  and as  $\text{Variance}(Y)-\hat{\sigma}$  increases.

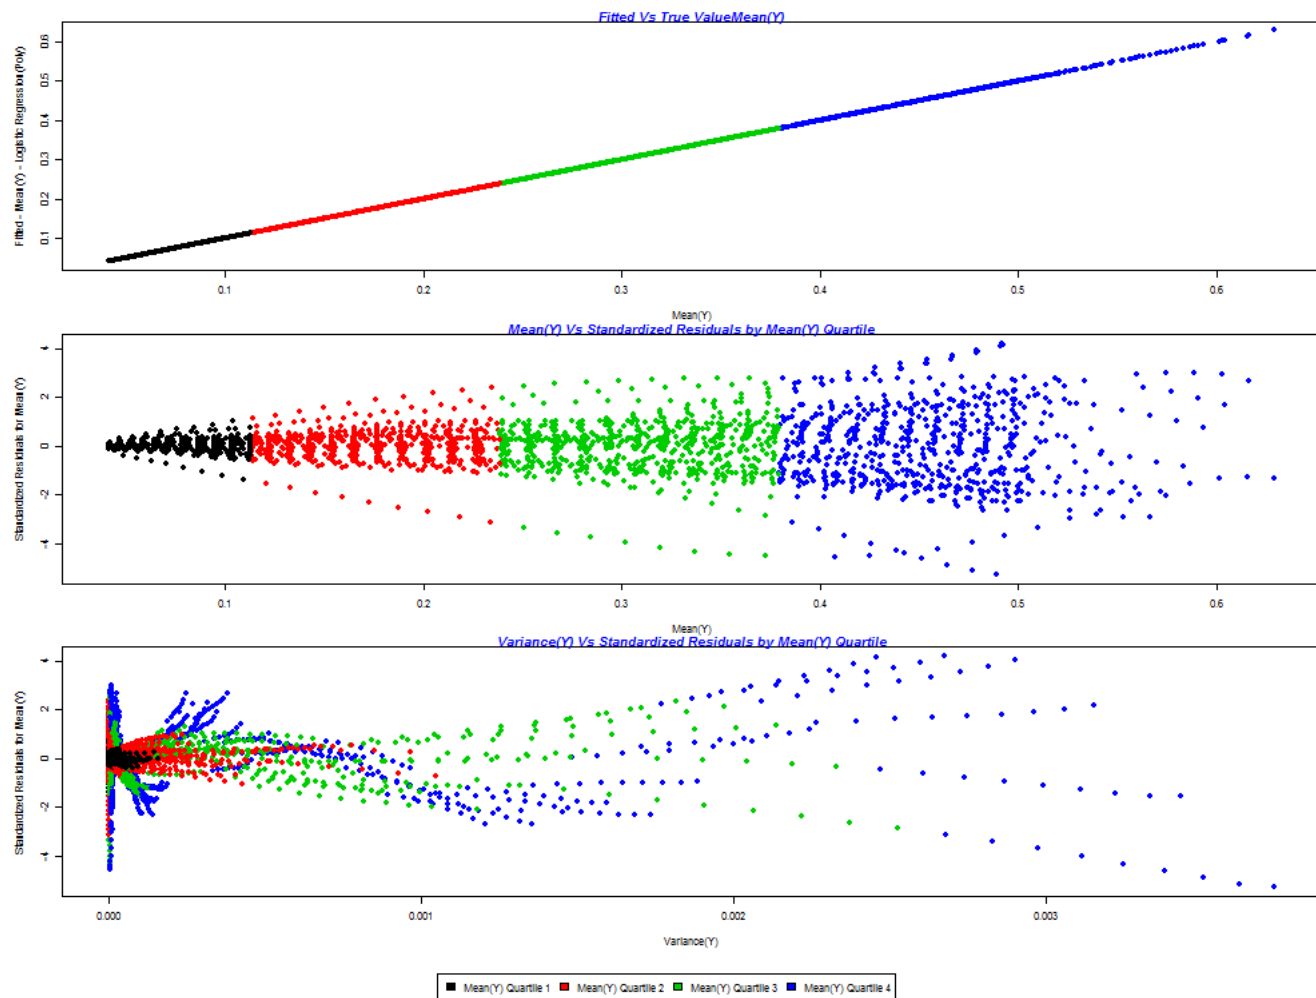


Figure 9: Logistic Regression Residuals for  $\text{Mean}(Y) - \hat{\mu}$  colored by  $\text{Mean}(Y) - \hat{\mu}$  Quartile



### 5.1.3 Linear Methods

#### Linear Regression

- The very low P-value for both the predictors shows that  $X_1, X_2$  are important for predicting the response  $\hat{\mu}$ . As expected the  $X_2$  has a positive coefficient that is approximately 20 times larger in magnitude than the negative coefficient of  $X_1$ . So  $X_2$  has a larger contribution to the response  $\hat{\mu}$ .

```
> summary(lm.mu.fit)
```

```
Call:
```

```
lm(formula = muhat ~ X1 + X2, data = data0.train)
```

```
Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.028385	-0.013942	-0.003463	0.008320	0.106545

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.0269571	0.0010017	26.91	<2e-16 ***
X1	-0.0065778	0.0001885	-34.90	<2e-16 ***
X2	0.1238886	0.0003272	378.69	<2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.02037 on 2762 degrees of freedom
```

```
Multiple R-squared:  0.9812,    Adjusted R-squared:  0.9812
```

```
F-statistic: 7.223e+04 on 2 and 2762 DF,  p-value: < 2.2e-16
```

- Figure (10) plots the fitted values for  $\text{Mean}(Y)-\hat{\mu}$  against the true values. The figure also plots the standardized residuals from the model with  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$ . The plots are colored to indicate the quartiles to which the predicted  $\text{Mean}(Y)-\hat{\mu}$  values belong. Ideally we should have a  $45^\circ$  line for the true and fitted values, the fit for this model seem moderate. The standardized residuals show a slight spread higher for upper quartiles of  $\text{Mean}(Y)-\hat{\mu}$  but remain constant as  $\text{Variance}(Y)-\hat{\sigma}$  increases.

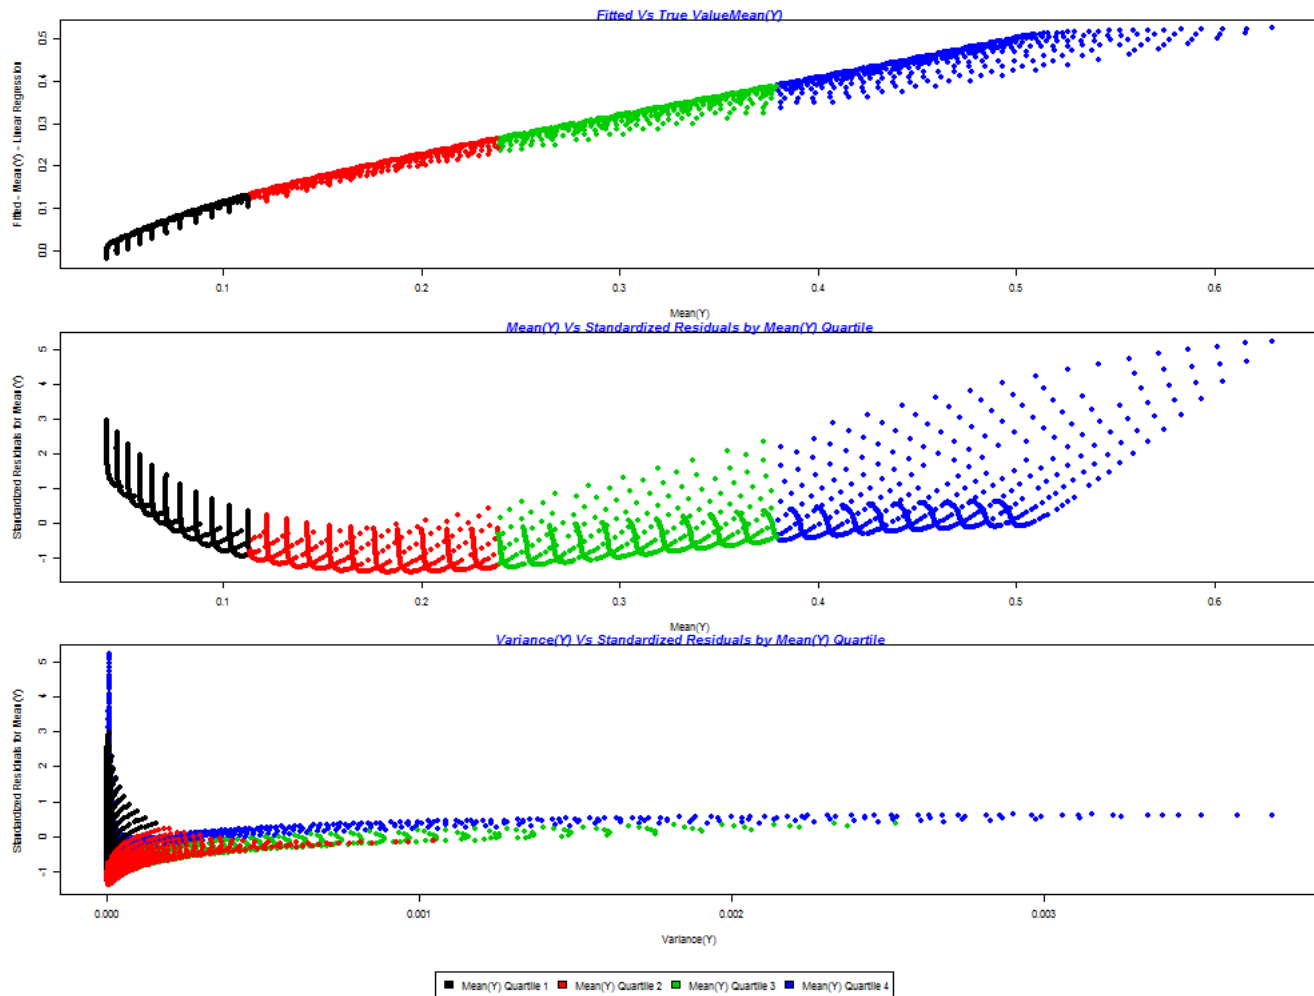


Figure 10: Linear Regression Residuals for  $\text{Mean}(Y) - \hat{\mu}$  colored by  $\text{Mean}(Y) - \hat{\mu}$  Quartile

## Linear Regression with Polynomial Terms

- The 10 fold cross validation to choose the optimum polynomial degree terms of  $X_1, X_1$  chooses degree 5 for  $X_1$  and degree 5 for  $X_2$ . The predictor terms with degree 1...4 of  $X_1$  and degree 1...3 of  $X_2$  have a very low P-value, so these predictors  $X_1, X_1^2, X_1^3, X_1^4, X_2, X_2^2, X_2^3$  are important for predicting the response  $\text{Mean}(Y) - \hat{\mu}$ . As expected the  $X_2$  has a positive coefficient that is approximately 10 times larger in magnitude than the next largest coefficient. So  $X_2$  has a larger contribution to the response  $\text{Mean}(Y) - \hat{\mu}$ . An  $R^2 = .99$  shows the model explains most of the variance seen in the response  $\text{Mean}(Y) - \hat{\mu}$ . The  $R^2 = .99$  for this model is slightly higher than the  $R^2 = .98$  for the previous discussed standard Linear Regression Model.

```
> summary(lmp.mu.fit)
```

Call:

```
lm(formula = muhat ~ poly(X1, poly.muhat.min.x1) + poly(X2, poly.muhat.min.x2),
    data = data0.train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.055339	-0.004731	0.000045	0.004683	0.053252

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.250342	0.000191	1310.527	< 2e-16 ***
poly(X1, poly.muhat.min.x1)1	-0.712481	0.010045	-70.927	< 2e-16 ***
poly(X1, poly.muhat.min.x1)2	0.481318	0.010045	47.916	< 2e-16 ***
poly(X1, poly.muhat.min.x1)3	-0.268686	0.010045	-26.748	< 2e-16 ***
poly(X1, poly.muhat.min.x1)4	0.103488	0.010046	10.301	< 2e-16 ***
poly(X1, poly.muhat.min.x1)5	-0.032485	0.010045	-3.234	0.00124 **
poly(X2, poly.muhat.min.x2)1	7.713699	0.010045	767.882	< 2e-16 ***
poly(X2, poly.muhat.min.x2)2	0.680636	0.010045	67.759	< 2e-16 ***
poly(X2, poly.muhat.min.x2)3	-0.295859	0.010046	-29.452	< 2e-16 ***
poly(X2, poly.muhat.min.x2)4	-0.010412	0.010045	-1.037	0.30003
poly(X2, poly.muhat.min.x2)5	0.022359	0.010045	2.226	0.02611 *

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 0.01004 on 2754 degrees of freedom

Multiple R-squared: 0.9955, Adjusted R-squared: 0.9954

F-statistic: 6.029e+04 on 10 and 2754 DF, p-value: < 2.2e-16

- Figure (11) plots the fitted values for  $\text{Mean}(Y)-\hat{\mu}$  against the true values. The figure also plots the standardized residuals from the model with  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$ . The plots are colored to indicate the quartiles to which the predicted  $\text{Mean}(Y)-\hat{\mu}$  values belong. Ideally we should have a  $45^\circ$  line for the true and fitted values, the fit for this model seems good. The standardized residuals show a slight spread higher for upper quartiles of  $\text{Mean}(Y)-\hat{\mu}$  but remain constant as  $\text{Variance}(Y)-\hat{\sigma}$  increases.

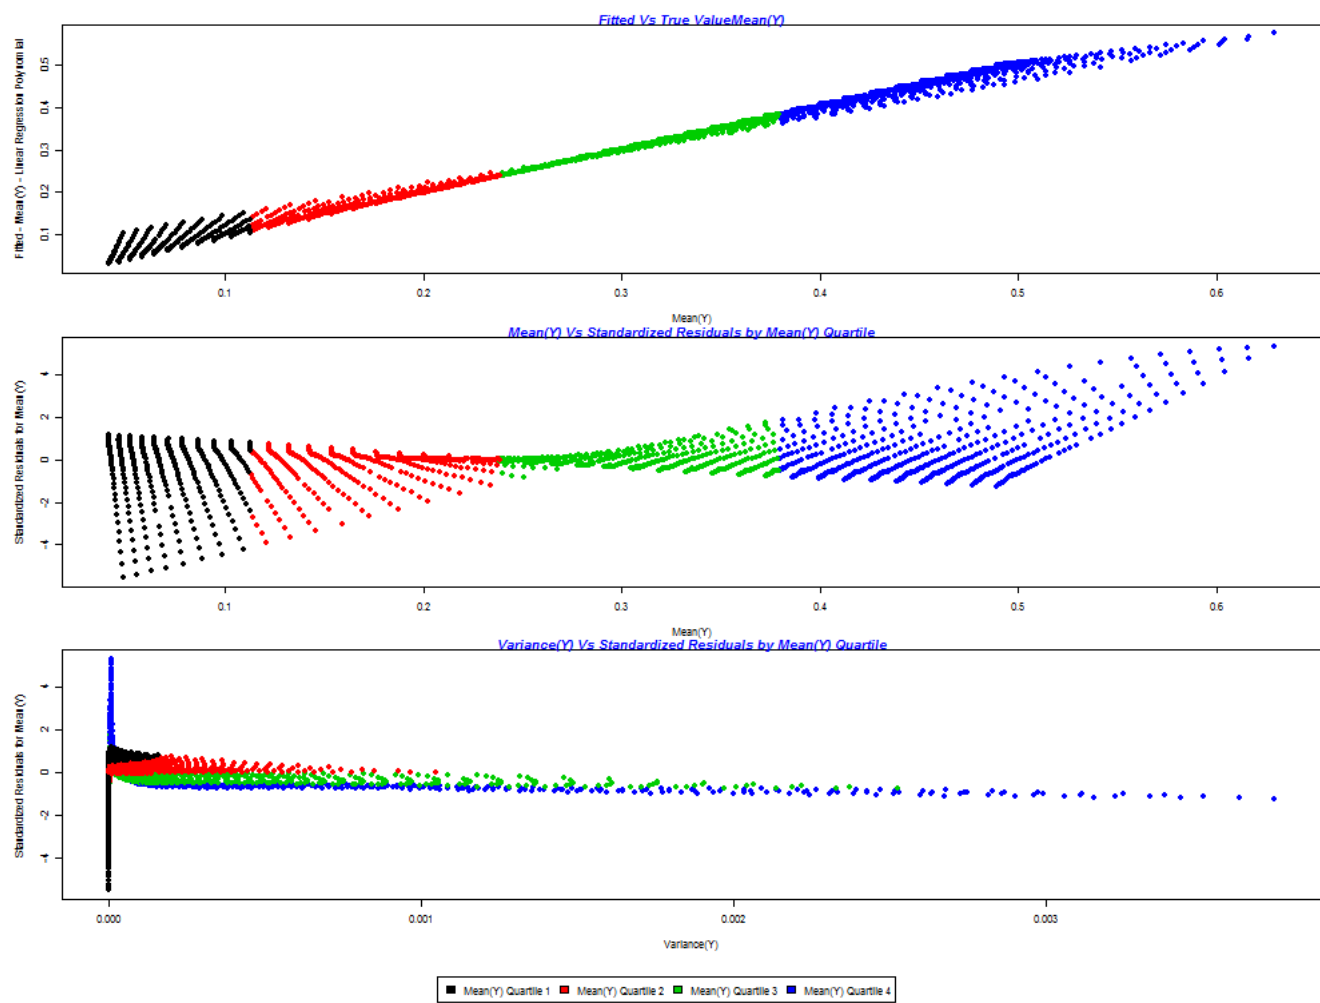


Figure 11: Linear Regression (Polynomial) Residuals for  $\text{Mean}(Y) - \hat{\mu}$  colored by  $\text{Mean}(Y) - \hat{\mu}$  Quartile

## LASSO with Polynomial Terms

- The 10 fold cross validation to choose the optimum polynomial degree terms of  $X_1, X_1$  for LASSO, along with a further ten fold validation to choose the best  $\lambda$  for these polynomial terms yields an value of  $\lambda = 1.429932e-05$ . The cross validation plot for LASSO plotting MSE for various values of  $\lambda$  is in figure (12) and shows the best  $\lambda$  value after which we see a sharp upward elbow in MSE.

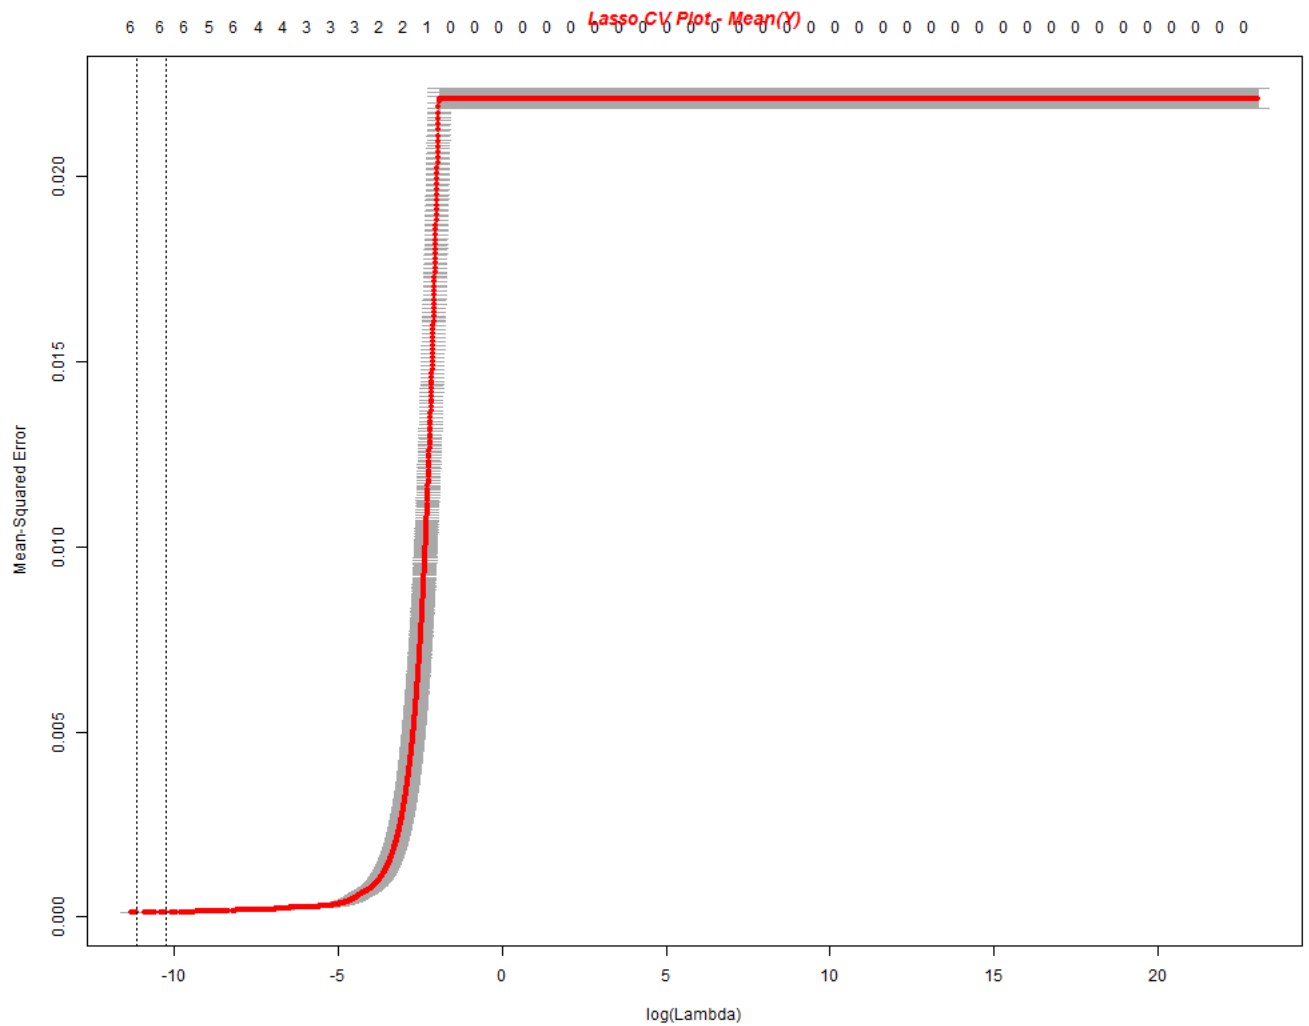


Figure 12: LASSO (Polynomial) Cross Validation Plot for  $\text{Mean}(\hat{Y}) - \hat{\mu}$  for choosing  $\lambda$

- The chosen  $\lambda$  shrinks the coefficients for several polynomial terms of  $X_1, X_2$ , choosing only  $X_1, X_1^2, X_1^3, X_2, X_2^2, X_2^4$ . However in contrast to the previously discussed Linear Regression with polynomial terms, the coefficients have been reduced in magnitude for  $X_1, X_2$  by a tenth. The coefficient of  $X_2$  is just slightly higher than the magnitude of the coefficient for  $X_1$  and  $X_2^2$  has a coefficient of almost equal in magnitude to  $X_1$ . The  $R^2 = .9953$  computed for this model is slightly higher than the  $R^2 = .98$  for the previous discussed standard Linear Regression Model.

```
> lasso.muhat.bestlam
[1] 1.429932e-05
>
>
> coef(lasso.mu.fit, s = "lambda.min")
8 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept) 0.0917578311
X1          -0.0384341029
X2           0.0529331096
X2.1         0.0078598984
X3          -0.0005229176
X2.2         0.0244319522
X3.1         .
X4.1        -0.0005124043
```

- Figure (13) plots the fitted values for  $\text{Mean}(Y) - \hat{\mu}$  against the true values. The figure also plots the standardized residuals from the model with  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$ . The plots are colored to indicate the quartiles to which the predicted  $\text{Mean}(Y) - \hat{\mu}$  values belong. Ideally we should have a  $45^\circ$  line for the true and fitted values, the fit for this model seems very good. The standardized residuals show a slight spread higher for upper quartiles of  $\text{Mean}(Y) - \hat{\mu}$  but remain constant as  $\text{Variance}(Y) - \hat{\sigma}$  increases.

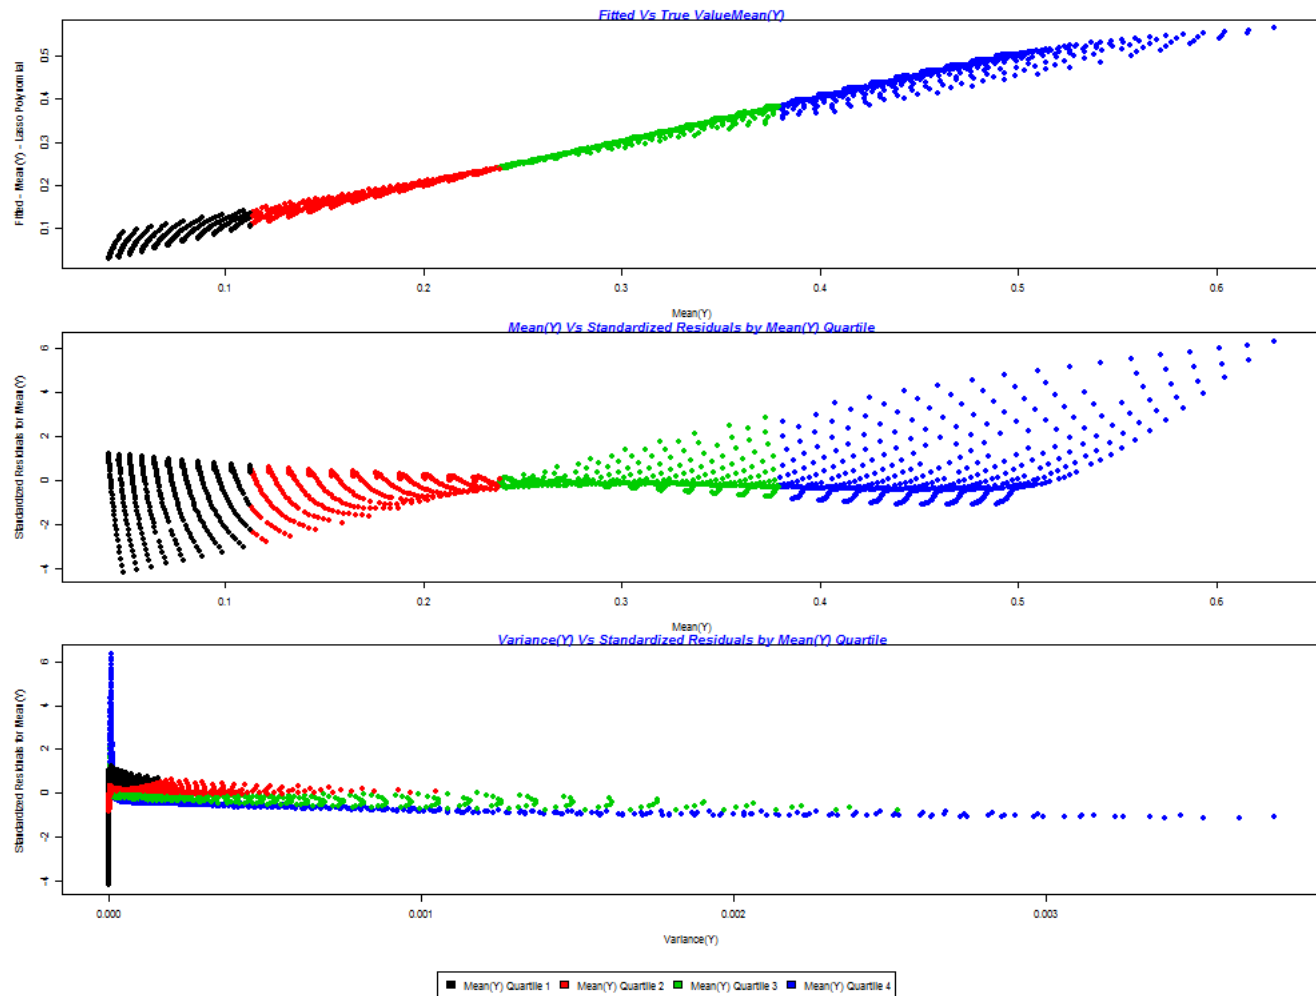


Figure 13: LASSO (Polynomial) Residuals for  $\text{Mean}(Y) - \hat{\mu}$  colored by  $\text{Mean}(Y) - \hat{\mu}$  Quartile



## 5.2 Predicting Variance of $Y - \hat{\sigma}$

### 5.2.1 Logistic Regression

- The P-Values of the coefficients are not very significant indicating that the predictors are not strong contributors to the response. The Residual deviance is low.

```
> summary(lr.v.fit)
```

Call:

```
glm(formula = Vhat ~ X1 * X2, family = binomial(logit), data = data0.train)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-0.0199292	-0.0021285	-0.0003975	0.0009398	0.0101573

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.524e+01	2.574e+01	-0.592	0.554
X1	8.855e-01	4.262e+00	0.208	0.835
X2	9.451e-01	8.055e+00	0.117	0.907
X1:X2	-4.239e-04	1.333e+00	0.000	1.000

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1.696732 on 2764 degrees of freedom  
Residual deviance: 0.043195 on 2761 degrees of freedom  
AIC: 9.2597

Number of Fisher Scoring iterations: 14

- Figure (14) plots the fitted values for  $\text{Variance}(Y) - \hat{\sigma}$  against the true values. The figure also plots the standardized residuals from the model with  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$ . The plots are colored to indicate the quartiles to which the predicted  $\text{Variance}(Y) - \hat{\sigma}$  values belong. Ideally we should have a  $45^\circ$  line for the true and fitted values, in this case plot appears to have a spread as the value of  $\text{Mean}(Y) - \hat{\mu}$  increases. The model appears to underestimate  $\text{Variance}(Y) - \hat{\sigma}$ . The standardized residuals also spread lower for upper quartiles of  $\text{Variance}(Y) - \hat{\sigma}$

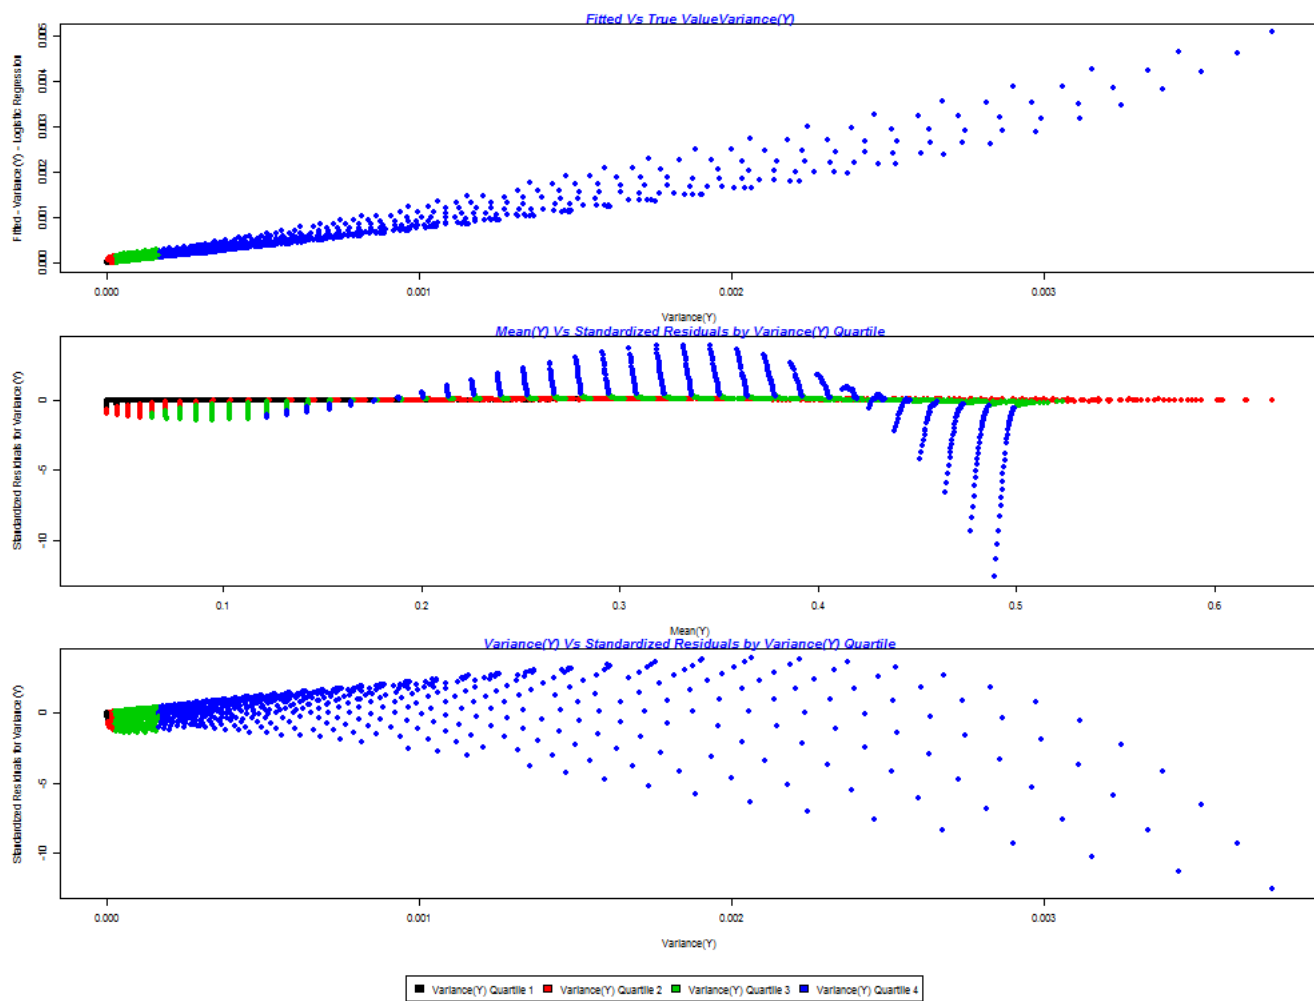


Figure 14: Logistic Regression Residuals for  $\text{Variance}(Y) - \hat{\sigma}$  colored by  $\text{Variance}(Y) - \hat{\sigma}$  Quartile

### 5.2.2 Logistic Regression with Polynomial Terms

- The 10 fold cross validation plot for MSE for various combinations of polynomial degree terms of  $X_1, X_2$  is shown in figure (15). The ten fold cross validation picks degree 4 for the interaction of  $X_1, X_2$  as the best model.

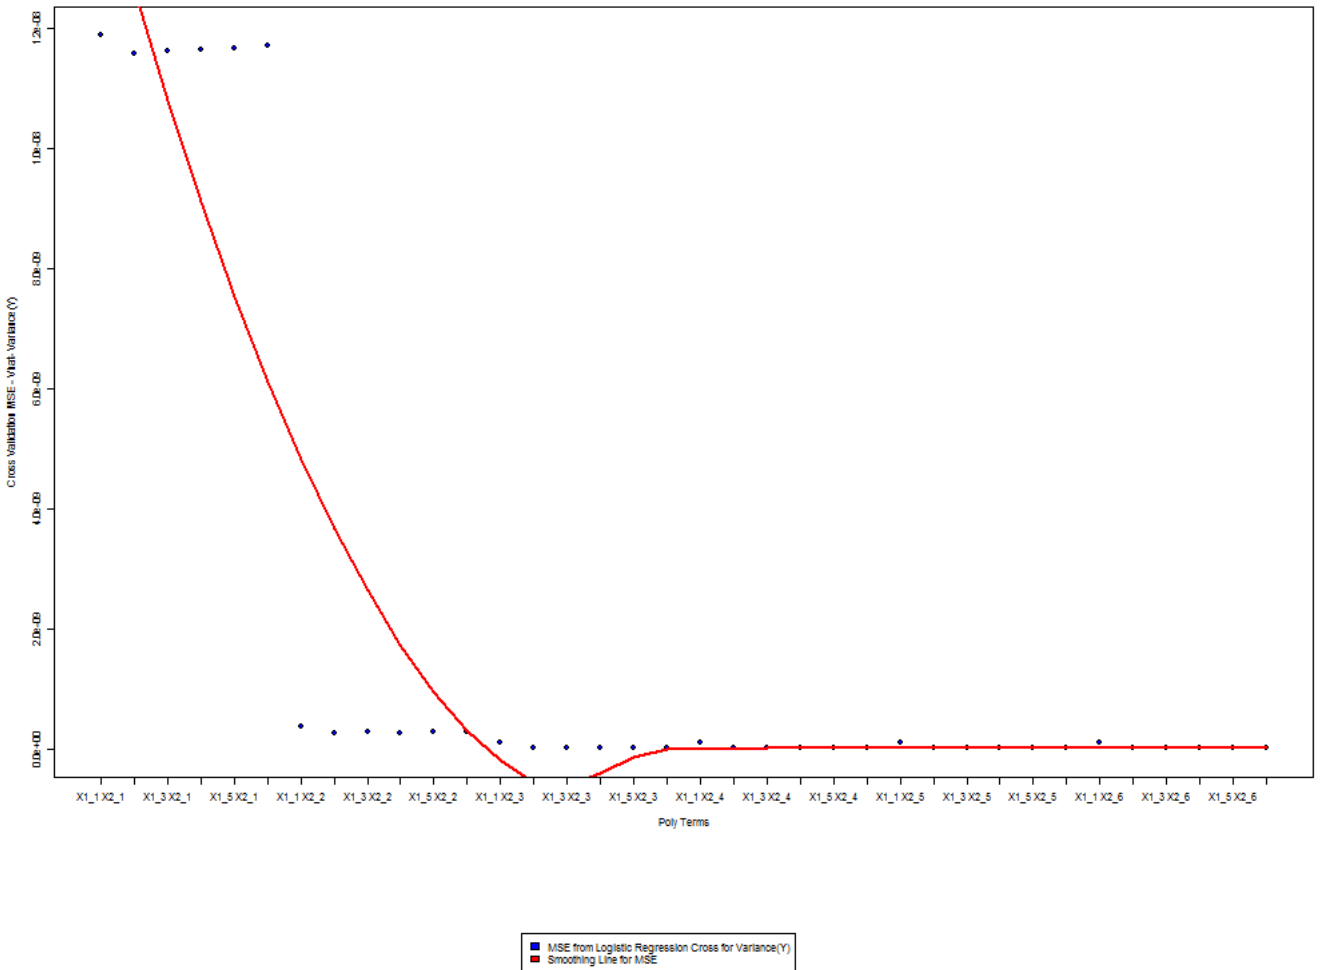


Figure 15: Cross Validation MSE for predicting Variance( $Y$ )- $\hat{\sigma}$  plotted against Predictor Combinations

- The P-values of the coefficients are not very significant, indicating that the any

single predictor itself is not a very strong contributor to the Variance(Y)- $\hat{\sigma}$  response in this model. The model also has a very low residual deviance of 5.1198e-05. This indicates a good fit.

```
> summary(lrp.v.fit)
```

Call:

```
glm(formula = Vhat ~ poly(X1, poly.lrp.vhat.min.x1) * poly(X2,
  poly.lrp.vhat.min.x2), family = binomial(logit), data = data0.train)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-6.831e-04	-5.645e-05	-1.440e-06	5.897e-05	4.655e-04

Coefficients:

	Estimate	Std. Err
(Intercept)	-1.069e+01	1.658e+01
poly(X1, poly.lrp.vhat.min.x1)1	9.506e+01	1.075e+01
poly(X1, poly.lrp.vhat.min.x1)2	2.842e+00	9.944e+00
poly(X1, poly.lrp.vhat.min.x1)3	-3.002e+00	8.527e+00
poly(X1, poly.lrp.vhat.min.x1)4	9.794e-01	4.727e+00
poly(X2, poly.lrp.vhat.min.x2)1	9.098e+01	1.218e+01
poly(X2, poly.lrp.vhat.min.x2)2	-2.918e+01	1.129e+01
poly(X2, poly.lrp.vhat.min.x2)3	4.658e+00	8.387e+00
poly(X2, poly.lrp.vhat.min.x2)4	-3.724e-01	4.616e+00
poly(X1, poly.lrp.vhat.min.x1)1:poly(X2, poly.lrp.vhat.min.x2)1	9.928e+01	7.904e+01
poly(X1, poly.lrp.vhat.min.x1)2:poly(X2, poly.lrp.vhat.min.x2)1	-1.051e+02	7.310e+01
poly(X1, poly.lrp.vhat.min.x1)3:poly(X2, poly.lrp.vhat.min.x2)1	1.416e+01	6.264e+01
poly(X1, poly.lrp.vhat.min.x1)4:poly(X2, poly.lrp.vhat.min.x2)1	2.414e+00	3.469e+01
poly(X1, poly.lrp.vhat.min.x1)1:poly(X2, poly.lrp.vhat.min.x2)2	3.491e+01	7.331e+01
poly(X1, poly.lrp.vhat.min.x1)2:poly(X2, poly.lrp.vhat.min.x2)2	-7.987e+00	6.774e+01
poly(X1, poly.lrp.vhat.min.x1)3:poly(X2, poly.lrp.vhat.min.x2)2	1.638e+01	5.799e+01
poly(X1, poly.lrp.vhat.min.x1)4:poly(X2, poly.lrp.vhat.min.x2)2	-7.784e+00	3.205e+01
poly(X1, poly.lrp.vhat.min.x1)1:poly(X2, poly.lrp.vhat.min.x2)3	-1.324e+01	5.468e+01
poly(X1, poly.lrp.vhat.min.x1)2:poly(X2, poly.lrp.vhat.min.x2)3	6.697e+00	5.054e+01
poly(X1, poly.lrp.vhat.min.x1)3:poly(X2, poly.lrp.vhat.min.x2)3	-6.783e-01	4.304e+01
poly(X1, poly.lrp.vhat.min.x1)4:poly(X2, poly.lrp.vhat.min.x2)3	6.267e-01	2.368e+01
poly(X1, poly.lrp.vhat.min.x1)1:poly(X2, poly.lrp.vhat.min.x2)4	1.583e+00	3.052e+01
poly(X1, poly.lrp.vhat.min.x1)2:poly(X2, poly.lrp.vhat.min.x2)4	-7.899e-01	2.817e+01
poly(X1, poly.lrp.vhat.min.x1)3:poly(X2, poly.lrp.vhat.min.x2)4	-7.512e-01	2.355e+01
poly(X1, poly.lrp.vhat.min.x1)4:poly(X2, poly.lrp.vhat.min.x2)4	1.268e-02	1.275e+01

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1.6967e+00 on 2764 degrees of freedom  
Residual deviance: 5.1198e-05 on 2740 degrees of freedom  
AIC: 51.26

Number of Fisher Scoring iterations: 18

- Figure (16) plots the fitted values for  $\text{Variance}(Y)-\hat{\sigma}$  against the true values. The figure also plots the standardized residuals from the model with  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$ . The plots are colored to indicate the quartiles to which the predicted  $\text{Variance}(Y)-\hat{\sigma}$  values belong. Ideally we should have a  $45^\circ$  line for the true and fitted values, in this case the line appears to be almost ideal, which is very good. The standardized residuals show a slight spread higher for upper quartiles of  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$

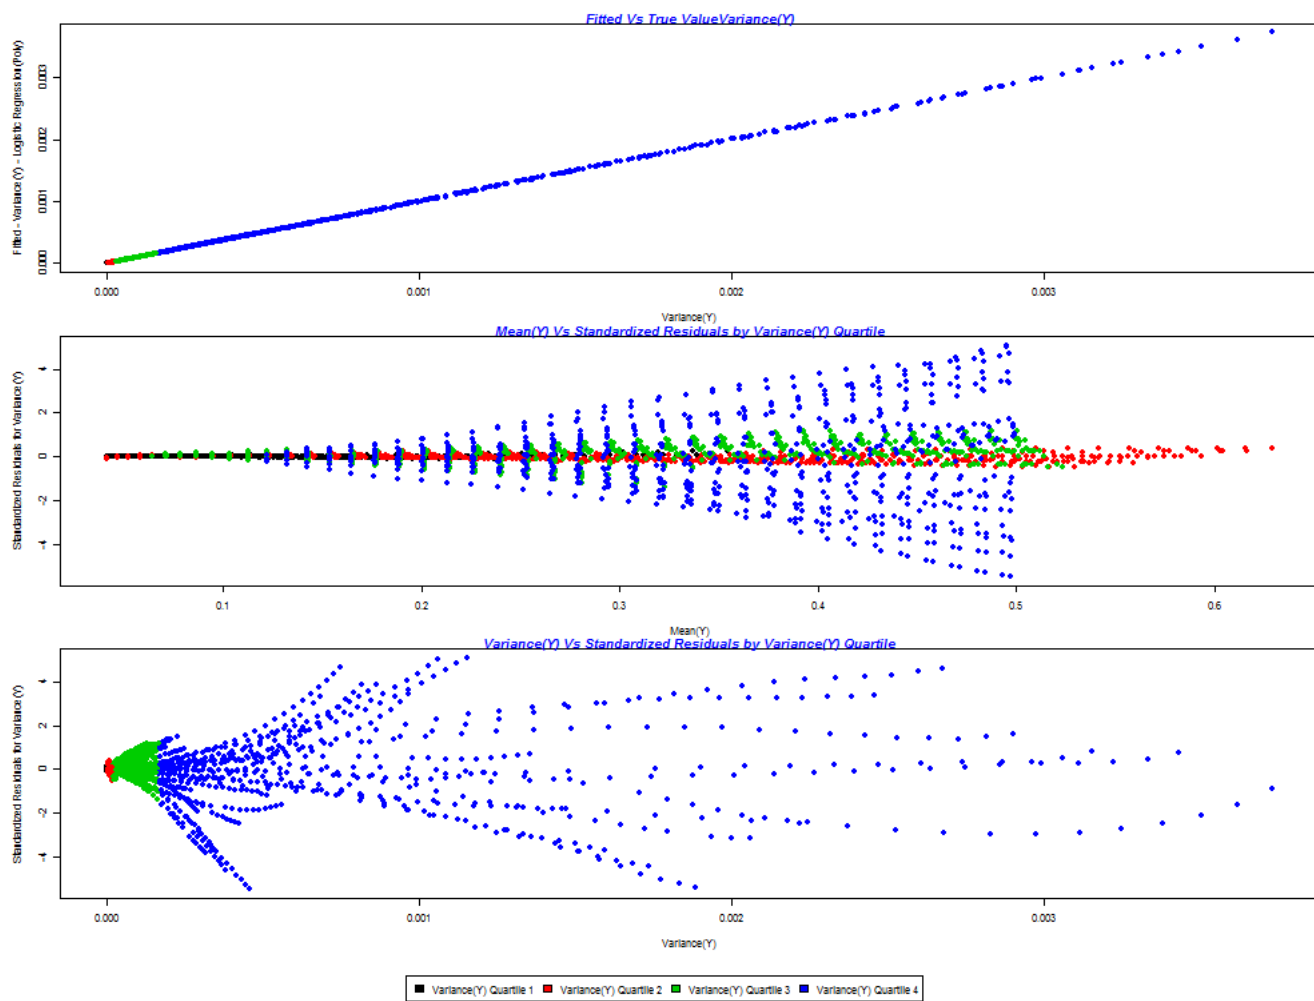


Figure 16: Logistic Regression Residuals for  $\text{Variance}(Y) - \hat{\sigma}$  colored by  $\text{Variance}(Y) - \hat{\sigma}$  Quartile

### 5.2.3 Linear Methods

#### Linear Regression

- The very low P-value for both the predictors shows that  $X_1, X_2$  are important for predicting the response  $\hat{\sigma}$ . As expected the  $X_1, X_2$  have positive coefficients that are comparable in magnitude. So both predictors have a approximately similar contribution to the response  $\hat{\sigma}$ . However the  $R^2 = 0.46$ , is very low and indicates that the model is does a poor job of explaining away the variance seen in  $\text{Variance}(Y) - \hat{\sigma}$

```
> summary(lm.v.fit)
```

```
Call:
```

```
lm(formula = What ~ X1 + X2, data = data0.train)
```

```
Residuals:
```

	Min	1Q	Median	3Q	Max
	-4.139e-04	-2.441e-04	-9.666e-05	1.318e-04	2.690e-03

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-5.808e-04	1.819e-05	-31.93	<2e-16 ***
X1	1.310e-04	3.423e-06	38.28	<2e-16 ***
X2	1.758e-04	5.942e-06	29.58	<2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.00037 on 2762 degrees of freedom
```

```
Multiple R-squared:  0.4603,    Adjusted R-squared:  0.4599
```

```
F-statistic: 1178 on 2 and 2762 DF,  p-value: < 2.2e-16
```

- Figure (17) plots the fitted values for  $\text{Variance}(Y) - \hat{\sigma}$  against the true values and shows a very poor fit. The plots for the standardized residuals from the model with  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$  show residuals increase with  $\text{Variance}(Y) - \hat{\sigma}$  and spread wider with  $\text{Mean}(Y) - \hat{\mu}$  for this model

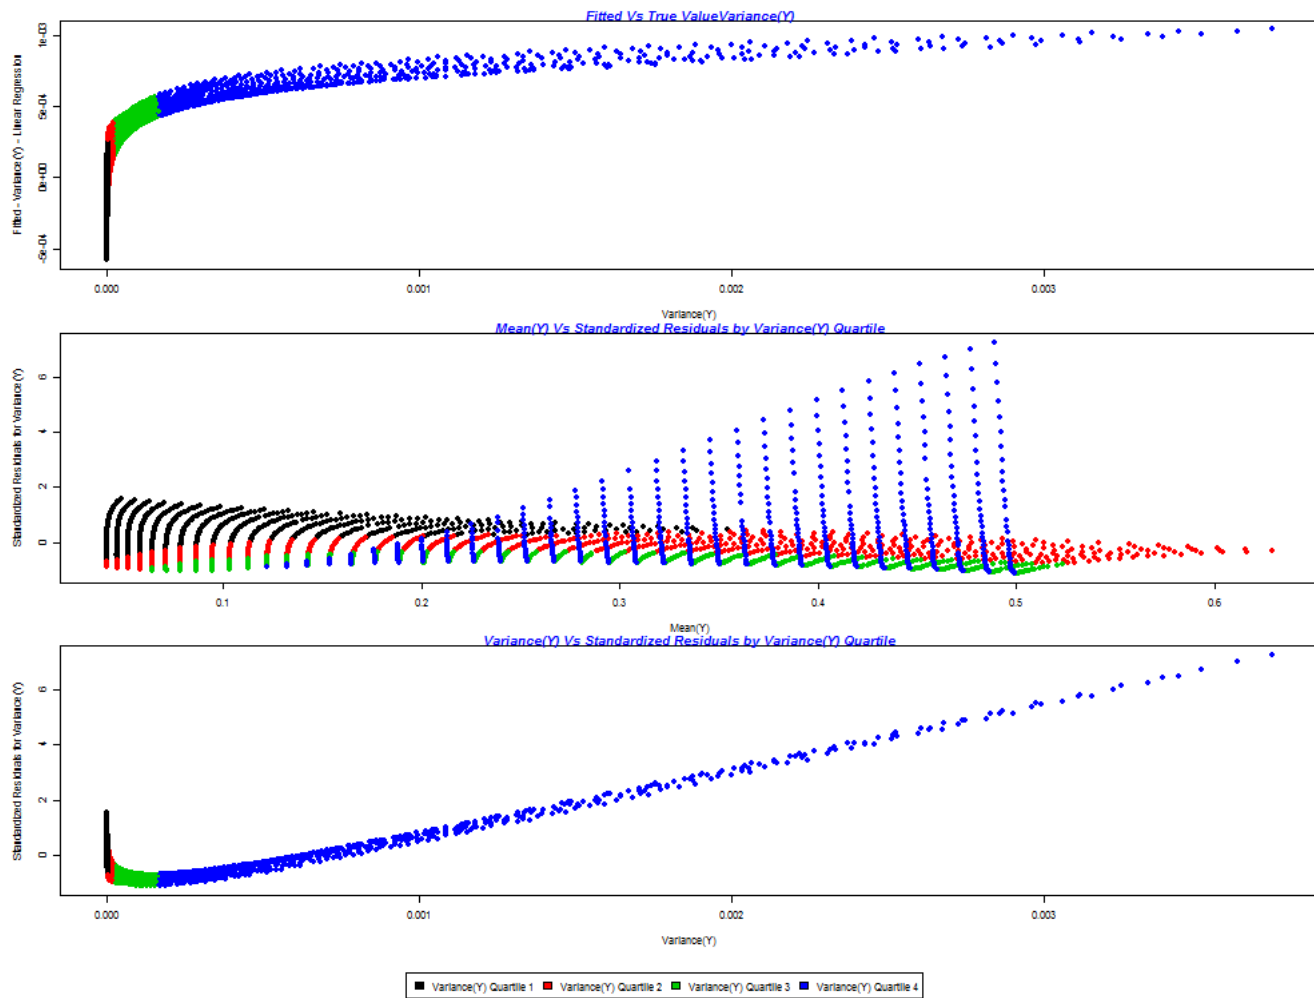


Figure 17: Linear Regression Residuals for  $\text{Variance}(Y) - \hat{\sigma}$  colored by  $\text{Variance}(Y) - \hat{\sigma}$  Quartile



## Linear Regression with Polynomial Terms

- The 10 fold cross validation to choose the optimum polynomial degree terms for  $X_1, X_1$  chooses degree 4 for  $X_1$  and degree 4 for  $X_2$ . The predictor terms with degree 1...3 of  $X_1$  and degree 1...2 of  $X_2$  have a very low P-value, so these predictors  $X_1, X_1^2, X_1^3, X_2, X_2^2$  are important for predicting the response  $\text{Variance}(Y)-\hat{\sigma}$ . As expected the terms in  $X_1, X_2$  have coefficients that are comparable in magnitude. So both predictors have a approximately similar contribution to the response  $\text{Variance}(Y)-\hat{\sigma}$ . An  $R^2 = .82$  shows the model explains only around 60% of the variance seen in the response  $\text{Variance}(Y)-\hat{\sigma}$ . The  $R^2 = .62$  for this model is slightly higher than the  $R^2 = .46$  for the previous discussed standard Linear Regression Model but still not very good.

```
> summary(lmp.v.fit)
```

Call:

```
lm(formula = Vhat ~ poly(X1, poly.Vhat.min.x1) + poly(X2, poly.Vhat.min.x2),
    data = data0.train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.066e-03	-1.616e-04	1.999e-05	1.544e-04	2.013e-03

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	2.276e-04	5.934e-06	38.358	< 2e-16 ***
poly(X1, poly.Vhat.min.x1)1	1.416e-02	3.121e-04	45.375	< 2e-16 ***
poly(X1, poly.Vhat.min.x1)2	9.092e-03	3.121e-04	29.134	< 2e-16 ***
poly(X1, poly.Vhat.min.x1)3	4.001e-03	3.121e-04	12.821	< 2e-16 ***
poly(X1, poly.Vhat.min.x1)4	1.357e-03	3.121e-04	4.349	1.42e-05 ***
poly(X2, poly.Vhat.min.x2)1	1.091e-02	3.121e-04	34.970	< 2e-16 ***
poly(X2, poly.Vhat.min.x2)2	2.899e-03	3.121e-04	9.290	< 2e-16 ***
poly(X2, poly.Vhat.min.x2)3	-6.763e-04	3.121e-04	-2.167	0.0303 *
poly(X2, poly.Vhat.min.x2)4	-5.100e-04	3.121e-04	-1.634	0.1023

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

Residual standard error: 0.0003121 on 2756 degrees of freedom

Multiple R-squared: 0.617, Adjusted R-squared: 0.6159

F-statistic: 555.1 on 8 and 2756 DF, p-value: < 2.2e-16

- Figure (18) plots the fitted values for  $\text{Variance}(Y) - \hat{\sigma}$  against the true values and shows a very poor fit. The plots for the standardized residuals from the model with  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$  show residuals increase with  $\text{Variance}(Y) - \hat{\sigma}$  and spread wider with  $\text{Mean}(Y) - \hat{\mu}$  for this model

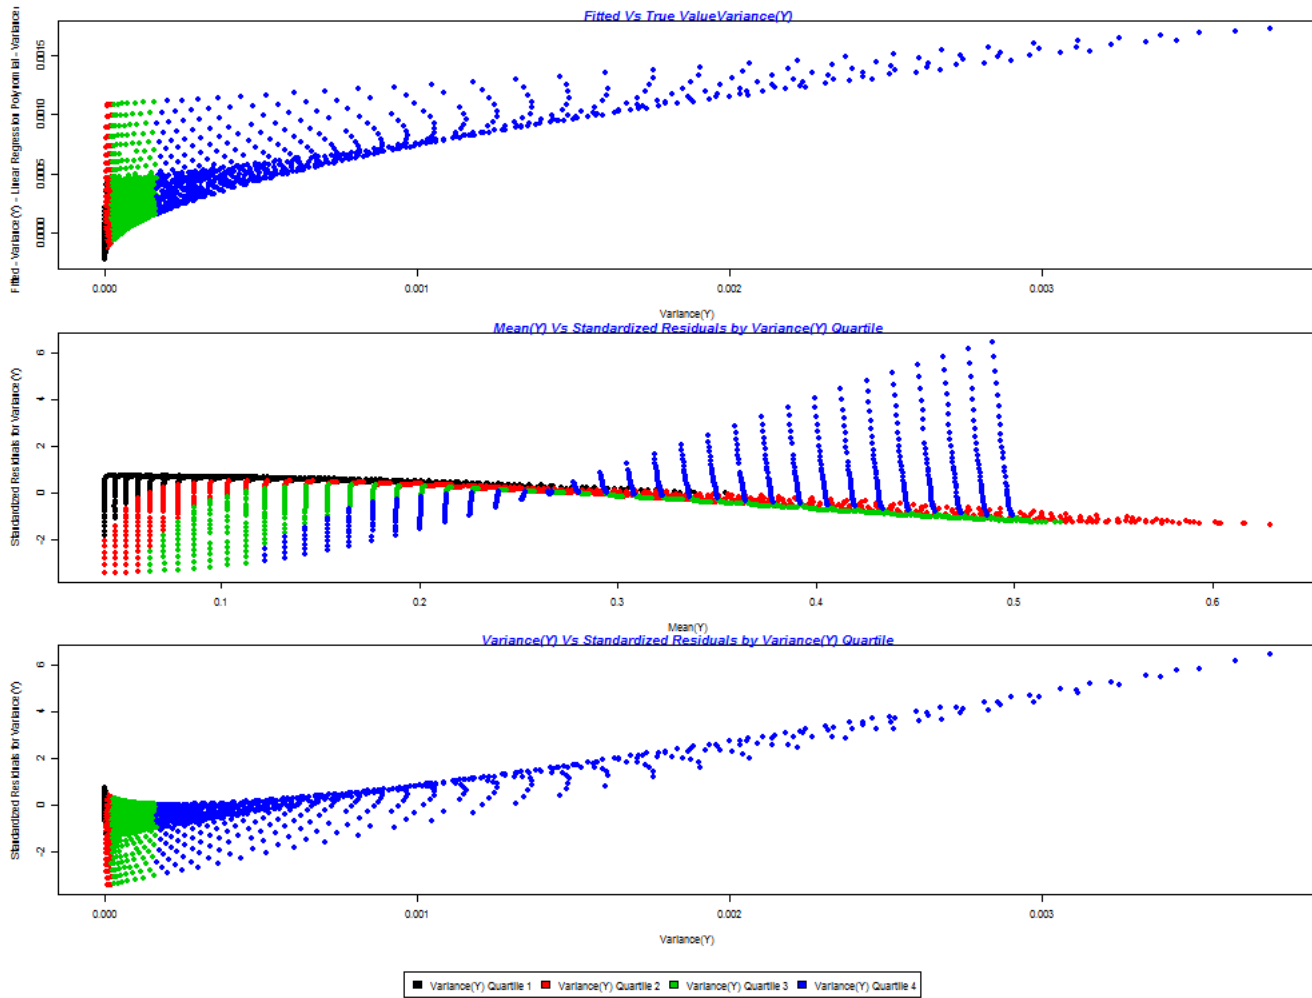


Figure 18: Linear Regression (Polynomial) Residuals for  $\text{Variance}(Y) - \hat{\sigma}$  colored by  $\text{Variance}(Y) - \hat{\sigma}$  Quartile

## LASSO with Polynomial Terms

- The 10 fold cross validation to choose the optimum polynomial degree terms of  $X_1, X_1$  for LASSO, along with a further ten fold validation to choose the best  $\lambda$  for these polynomial terms yields an value of  $\lambda = 1.429932e-05$ . The cross validation plot for LASSO plotting MSE for various values of  $\lambda$  is in figure (19) and shows the best  $\lambda$  value after which we see a sharp upward elbow in MSE.

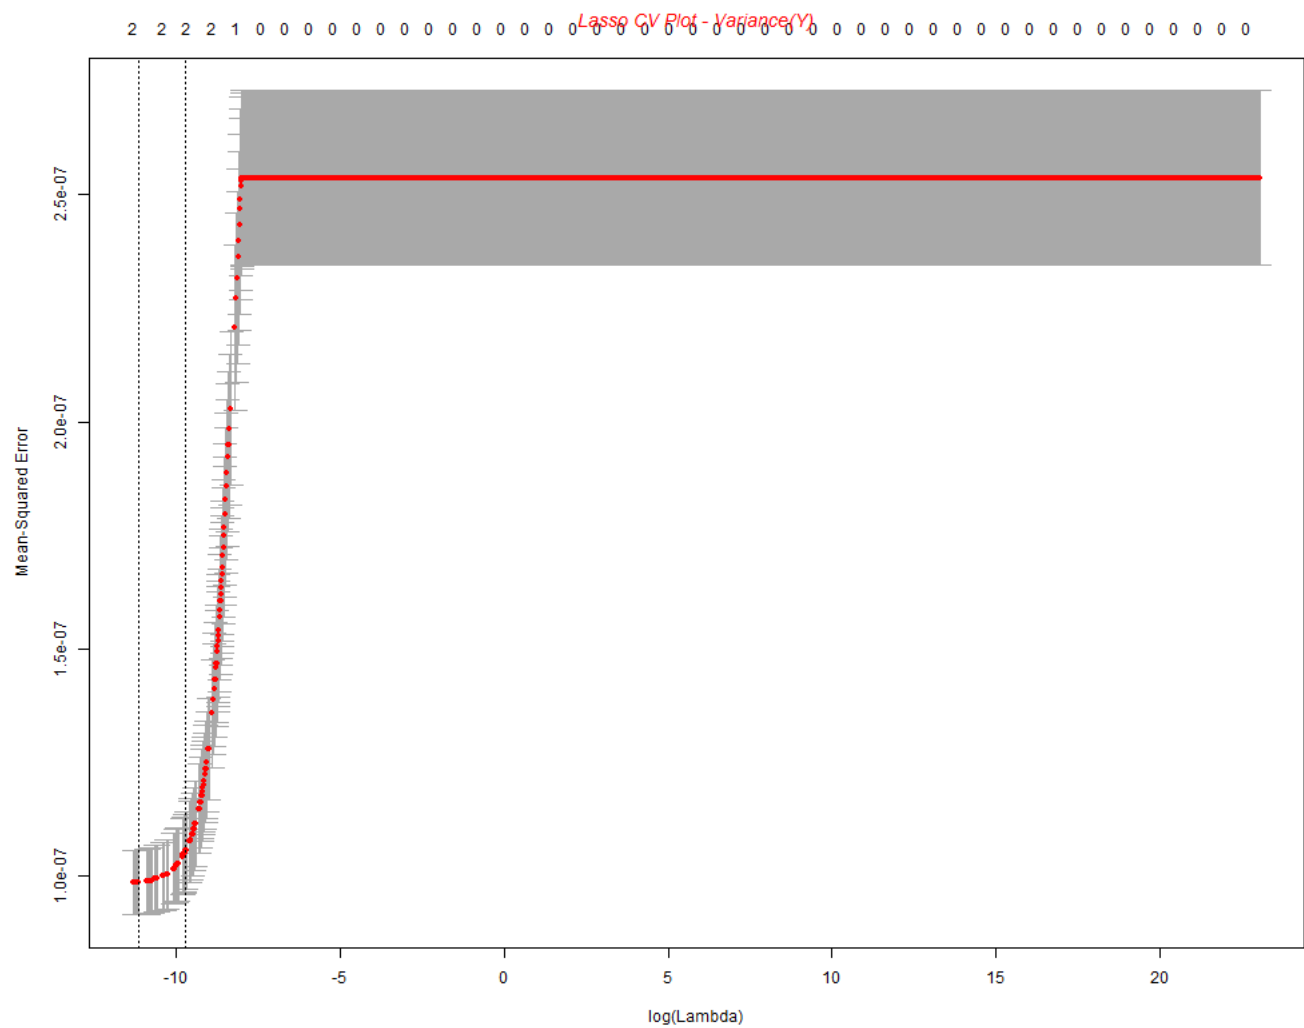


Figure 19: LASSO (Polynomial) Cross Validation Plot for  $\text{Variance}(Y) - \hat{\sigma}$  for choosing  $\lambda$

- The chosen  $\lambda$  shrinks the coefficients for several polynomial terms of  $X_1, X_2$ , choosing only  $X_1^5, X_2^2$ . However in contrast to the previously discussed Linear Regression with polynomial terms, The coefficients have been vastly reduced in magnitude. The coefficient of  $X_2^2$  is just higher than the magnitude of the coefficient for  $X_1^5$ . The  $R^2 = .61$  for this model is only slightly higher than the  $R^2 = .46$  for the previous discussed standard Linear Regression Model.

```
> lasso.Vhat.bestlam
[1] 1.429932e-05
```

```
> coef(lasso.v.fit, s = "lambda.min")
8 x 1 sparse Matrix of class "dgCMatrix"
              1
(Intercept) -2.022018e-04
X1           .
X2           .
X2.1         .
X3           .
X4           .
X5           7.202326e-08
X2.2         4.118209e-05

> print(lasso.Vhat.R.squared)
[1] 0.61413
```

- Figure (20) the plot of the fitted values for  $\text{Variance}(Y) - \hat{\sigma}$  against the true values again shows a very poor fit. The plots for the standardized residuals from the model with  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$  show residuals increase with  $\text{Variance}(Y) - \hat{\sigma}$  and spread wider with  $\text{Mean}(Y) - \hat{\mu}$  for this model

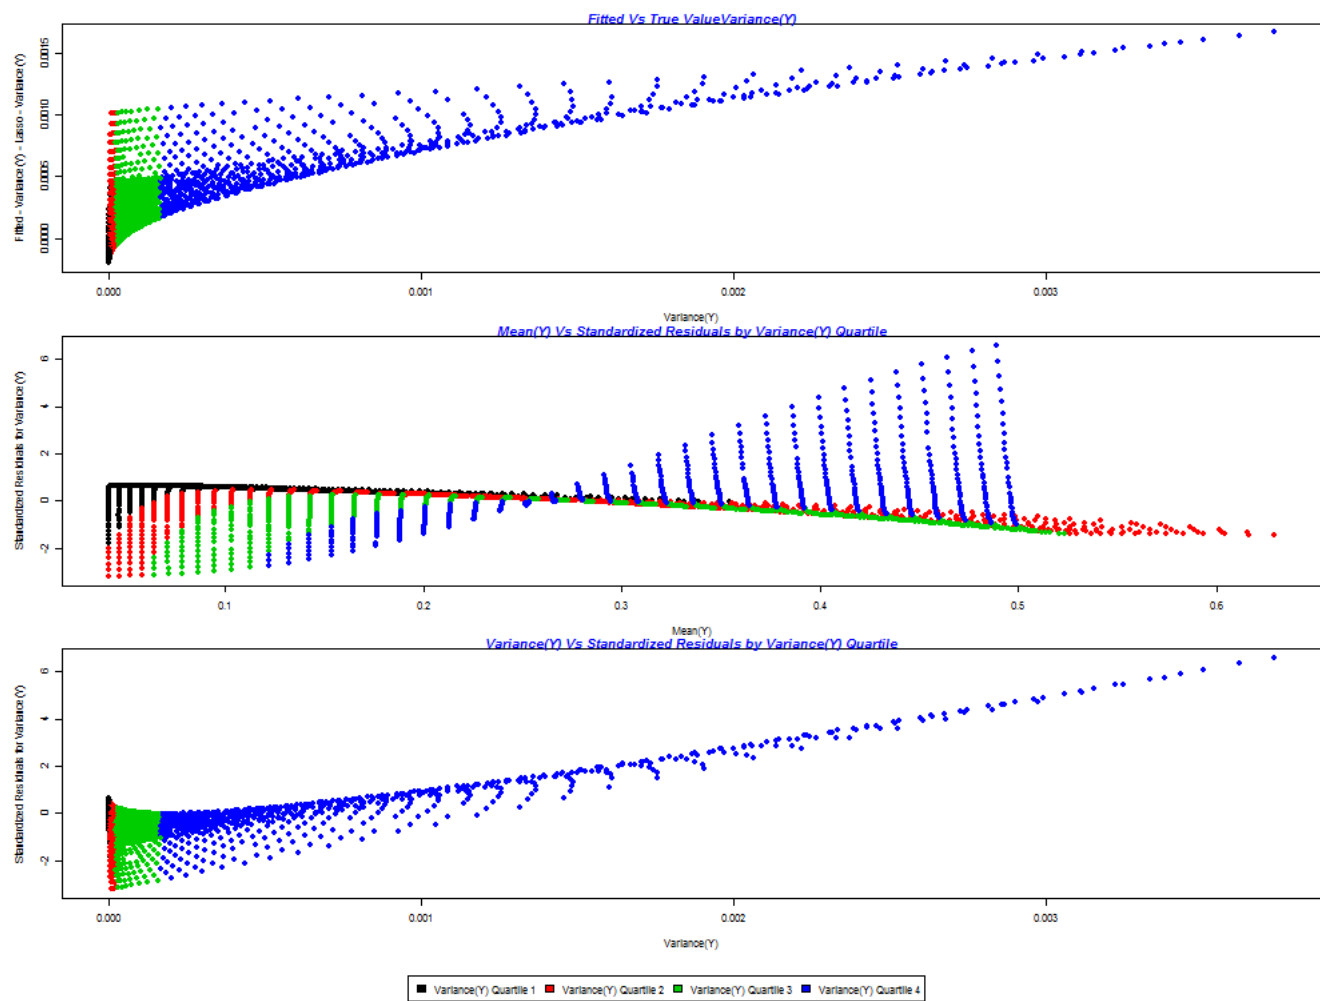


Figure 20: LASSO (Polynomial) Residuals for  $\text{Variance}(Y) - \hat{\sigma}$  colored by  $\text{Variance}(Y) - \hat{\sigma}$  Quartile

## 6 Test Results

### 6.1 Cross Validation Results

- The following table (2) tabulates the optimum parameters and predictors chosen for each of the models for predicting  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$ . These optimum parameters are picked after 10 fold cross validation from a pool of various candidate models. The models are chosen based on the lowest average MSE.

	Mean(Y)					Variance(Y)				
	Logistic Regression(LR)	LR (Polynomial)	Linear Regression (LM)	LM (Polynomial)	LASSO(Poly)	LR	LR(Poly)	LM	LM(Poly)	LASSO(Poly)
Predictors Used	$X_1$ $X_2$	$X_1 \dots X_1^6$ $X_2 \dots X_2^6$	$X_1$ $X_2$	$\text{Poly}(X_1, 5)$ $\text{Poly}(X_2, 5)$	$X_1, X_1^2, X_1^3$ $X_2, X_2^2, X_2^3$	$X_1$ $X_2$	$X_1 \dots X_1^4$ $X_2 \dots X_2^4$	$X_1$ $X_2$	$\text{Poly}(X_1, 4)$ $\text{Poly}(X_2, 4)$	$X_1^2$ $X_2^2$
Parameters Used					$\lambda = 1.43e-05$					$\lambda = 1.43e-05$

Table 2: Parameters and Predictors chosen for models by 10 fold cross validation

### 6.2 Bootstrapping 100 Iterations Test Results

- Tables (3) and (4) tabulate the bootstrapping results from 100 iterations for both  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$  predictions. Bootstrapping is performed using the optimum models chosen by cross validation. The table shows Average MSE and the variance in MSE for each model across 100 bootstrapping iterations. The results are for models predicting both the  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$ .

	Logistic Regression	Logistic Regression with Polynomial Terms	Linear Regression	Linear Regression with Polynomial terms	LASSO(Pol)
Average MSE(Mean(Y))	0.0003824652	1.017951e-08	0.0004166204	0.00010101	0.0001077252
Variance MSE(Mean(Y))	3.720936e-10	4.794325e-19	6.280023e-10	5.819501e-11	8.059884e-11

Table 3: Bootstrapping results for the models evaluated for predicting  $\text{Mean}(Y)-\hat{\mu}$

	Logistic Regression	Logistic Regression with Poly Terms	Linear Regression	Linear Regression with Poly terms	LASSO(Poly)
Average MSE(Variance(Y))	1.12547e-08	7.419421e-12	1.330242e-07	9.554755e-08	9.58125e-08
Variance MSE(Variance(Y))	4.471789e-18	4.082094e-25	1.376597e-16	5.702572e-17	6.542574e-17

Table 4: Bootstrapping results for the models evaluated for predicting  $\text{Variance}(Y)-\hat{\sigma}$

- Figure (21) box plots the MSE from each of the 100 bootstrap iterations for both the  $\text{Mean}(Y)-\hat{\mu}$  and  $\text{Variance}(Y)-\hat{\sigma}$  predictor models. The box plot shows the variance from each model.

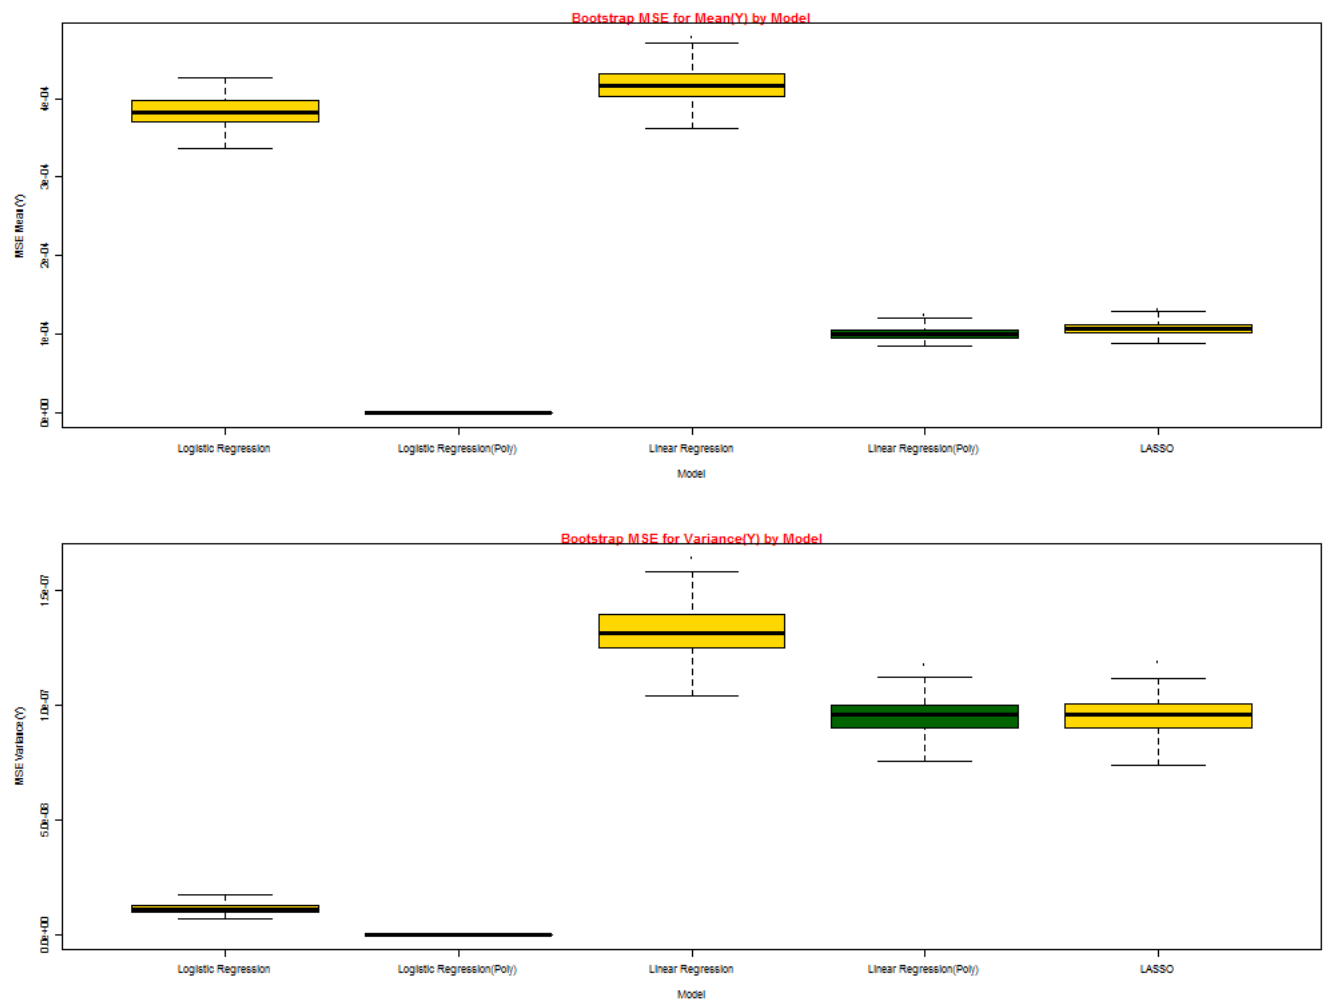


Figure 21: Box-plot of Bootstrapping MSE results of both  $\text{Mean}(Y) - \hat{\mu}$  and  $\text{Variance}(Y) - \hat{\sigma}$  predictor models

- Figure (22) summarizes the bootstrap results for each model predicting  $\text{Mean}(Y) - \hat{\mu}$  by plotting the average MSE and Variance in MSE.

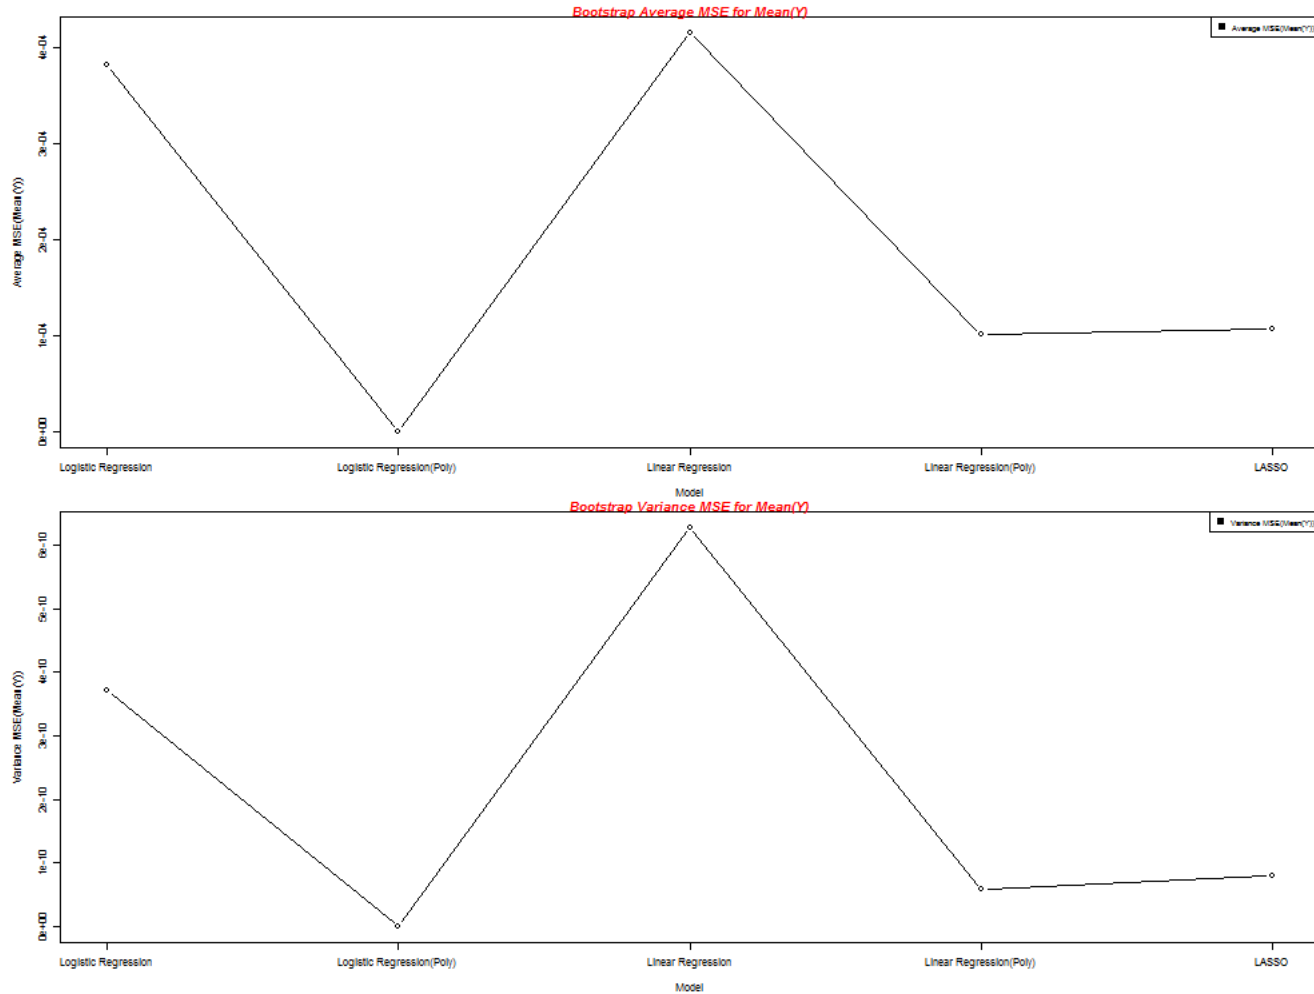


Figure 22: Bootstrapping 100 Iterations - Average MSE , Variance MSE for the  $\text{Mean}(Y) - \hat{\mu}$  prediction models

- Figure (23) summarizes the bootstrap results for each model predicting  $\text{Variance}(Y) - \hat{\sigma}$  by plotting the average MSE and Variance in MSE.



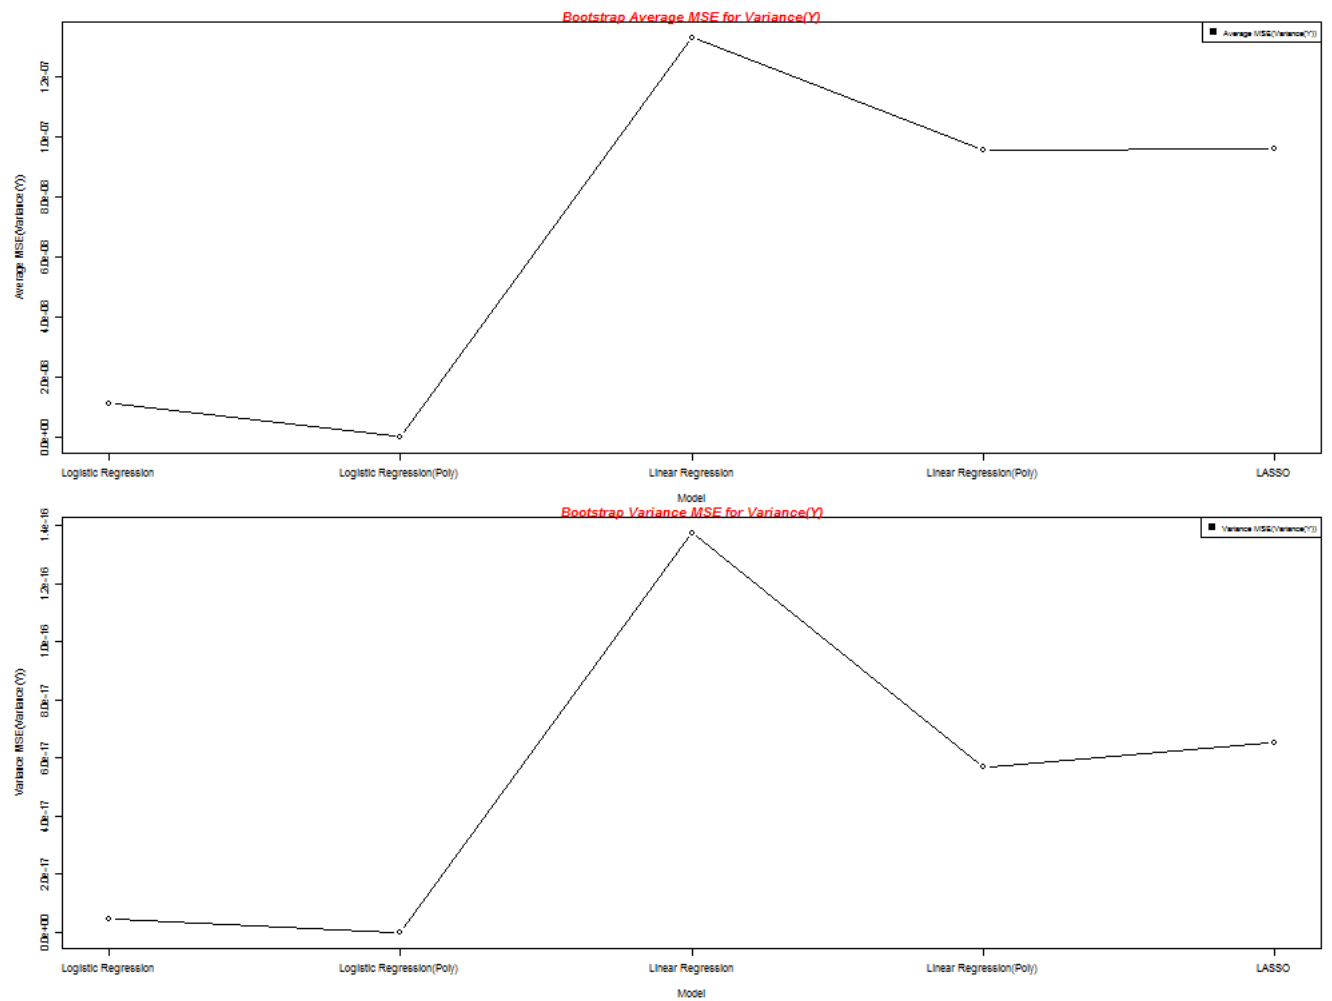


Figure 23: Bootstrapping 100 Iterations - Average MSE , Variance MSE for the  $\text{Variance}(Y)-\hat{\sigma}$  prediction models

### 6.3 T-test and Wilcox Test for Mean(Y)- $\hat{\mu}$ prediction models

- From the bootstrapping MSE results for the Mean(Y)- $\hat{\mu}$  predictor methods in table (3) and the plot in (22) The polynomial Logistic Regression Model has the lowest average MSE and variance in MSE.
- We run a T Test and a Wilcox test on these bootstrapping MSE with the null hypothesis  $H_0$  that these models are statistically similar at the  $\alpha = 5\%$  significance. The MSE from Polynomial Logistic Regression is compared to the MSE from other two models in a paired test.
- The P values from the T Test and Wilcox test results for the Mean(Y)- $\hat{\mu}$  predictor methods compared to Polynomial Logistic Regression are in table (5). The P-values are extremely low ( $<< .025$ ). This means that if null hypothesis  $H_0$  where true then the chances of seeing these results from the other models is extremely low. But since we are seeing these results, we can reject the null hypothesis  $H_0$  at the  $\alpha = 5\%$  significance and accept that alternate hypothesis, Polynomial Logistic Regression is the model with lowest MSE for predicting Mean(Y)- $\hat{\mu}$

	Logistic Regression	Linear Regression	Logistic Regression (Polynomial)	Lasso
T-Test	1.602735e-130	5.707506e-123	3.165143e-113	5.107431e-109
Wilcox Test	3.955912e-18	3.955912e-18	3.955912e-18	3.955912e-18

Table 5: T-test and Wilcox Test on Bootstrapping Results - P-Value Mean(Y)- $\hat{\mu}$  for models Vs Polynomial Logistic Regression

### 6.4 T-test and Wilcox Test for Variance(Y)- $\hat{\sigma}$ prediction models

- From the bootstrapping MSE results for the Variance(Y)- $\hat{\sigma}$  predictor models in table (3) and the plot in (23), the polynomial Logistic Regression Model has the lowest average MSE and variance in MSE with the LASSO model a close.
- We run a T Test and a Wilcox test on these bootstrapping MSE with the null hypothesis  $H_0$  that these models are statistically similar at the  $\alpha = 5\%$  significance. The MSE from Polynomial Logistic Regression is compared to the MSE from other two models in a paired test.
- The P values from the T Test and Wilcox test results for the Variance(Y)- $\hat{\sigma}$  predictor methods compared to Polynomial Linear Regression are in table (6). The

P-values are extremely low ( $<< .025$ ) for all methods. This means that if null hypothesis  $H_0$  where true then the chances of seeing these results from the standard Linear Regression model is extremely low. But since we are seeing these results, we can reject the null hypothesis  $H_0$  at the  $\alpha = 5\%$  significance and accept the alternate hypothesis that the Polynomial Logistic Regression is the model with lowest MSE for predicting  $\text{Variance}(Y) - \hat{\sigma}$ .

	Logistic Regression	Linear Regression	Linear Regression(Polynomial)	Lasso
T-Test	1.249724e-74	1.341322e-106	2.769225e-111	1.814889e-108
Wilcox Test	3.955912e-18	3.955912e-18	3.955912e-18	3.955912e-18

Table 6: T-test and Wilcox Test on Bootstrapping Results - P-Value  $\text{Variance}(Y) - \hat{\sigma}$  for models Vs Polynomial Logistic Regression

## 7 Conclusions

### 7.1 Mean(Y) - $\hat{\mu}$

- From the bootstrap results table (3) and plot (22) we can see that the Logistic Regression model with polynomial terms is the best performing model for predicting  $\text{Mean}(Y) - \hat{\mu}$ . The results from the T test and Wilcox test on the bootstrapping results in table (5) confirm that the MSE results for Polynomial Logistic Regression model are the best. As assumed earlier this makes sense as variance of  $Y$  is heteroscedastic and changes with  $X_1, X_2$  and consequently with  $\text{Mean}(Y) - \hat{\mu}$  as seen in the plot (5), Logistic regression does not assume homoscedasticity while the OLS based Linear Regression models do. Also the value of  $Y$  is in the range of  $0 \dots 1$ , so it can be thought of as  $P(X)$  which can be predicted by the Logistic Regression Model. Furthermore a plot of  $\text{Mean}(Y)$  with  $X_2$  in (4) takes the form of a sigmoid function, which is the model of Logistic Regression.
- Between the Logistic Regression standard model and the Logistic Regression Polynomial model, the Logistic Regression Polynomial model clearly outperforms the former based on the results seen in bootstrapping. While the polynomial terms might tend to over fit, this choice is made based on bootstrapping and cross validation MSE on a holdout test set every time. So predictions of  $\text{Mean}(Y) - \hat{\mu}$  for the test data provided will be made using this Polynomial Logistic Regression model.

## 7.2 Variance(Y) - $\hat{\sigma}$

- From the bootstrap results table (4) and plot (23) we can see that the Logistic Regression model with polynomial terms is again the best performing model for predicting Variance(Y)- $\hat{\sigma}$ . The results from the T test and Wilcox test on the bootstrapping results in table (6) confirm that the MSE results for Logistic Regression model with polynomial terms are the best.
- There is another reason for choosing the Logistic Regression Model. This model ensures that the values  $P(X)$  are in the range  $0 \dots 1$ . We know that Variance(Y) -  $\hat{\sigma}$  has to be in the range  $0 \dots 1$ , since Y is also in range  $0 \dots 1$ . So Variance(Y) can be modeled as  $P(X)$  of the Logistic Regression Model. Since no assumption is made of the distribution this fits well with Logistic Regression.
- Furthermore a plot of Variance(Y) -  $\hat{\sigma}$  with  $X_2$  and  $X_1$  in (4) resembles the form of a sigmoid function, which is the model of Logistic Regression.
- Between the Logistic Regression standard model and the Logistic Regression Polynomial model, the Logistic Regression with Polynomial terms model clearly outperforms the former in the bootstrap test results. While the polynomial terms might tend to over fit, this choice of this model is made based on bootstrapping and cross validation on hold out test sets each time. So predictions of Variance(Y)- $\hat{\sigma}$  for the test data provided will be made using this Logistic Regression with Polynomial terms model.

## 8 Appendix

The Following is the R source code used for this project. It is also available in file `ajdsouza31-midterm.R`

```
#  
#ajdsouza31 - ISYE7406Q - MidTerm  
#  
##### Some R codes for take-home midterm of ISyE 7406  
#####  
  
set.seed(20160227) ### set the random seed  
  
library(lattice)
```

```

library(glmnet)
library(corrplot)
library(GGally)

#-----
# Functions
#-----

#-----
# X1 vs X2 with display data
#-----
scatterplot.xy <- function(file.name,
data.xy,
display.data,
title.label,
title.main ) {

png(paste(file.name,"png",sep="."),width=1000,height=800)

m <- matrix(c(1,2,2),nrow = 3,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(0.80,0.1,.1))

par(mar = c(4,4,1,0))

plot(data.xy$X1,data.xy$X2,type='p', pch=21, col=as.numeric(display.data),
bg=as.numeric(display.data),
xlab="X1", ylab="X2")

title(main=title.main, col.main='red', font.main=4,outer=FALSE)

par(mar = c(0,0,0,0))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")

legend(x="bottom", inset=0, paste(title.label,levels(as.factor(Vqt))),
fill=levels(as.factor(display.data)) ,cex=1, horiz = TRUE)

dev.off()

```

```

}

#-----
# X vs Residuals, by display data
#-----
plotxy.residuals <- function (
  file.name,
  data,
  model.residuals,
  std.residuals,
  display.data,
  true.value,
  fitted.value,
  title.fitted,
  title.legend,
  title.model
) {

  png(paste(file.name,"png",sep="."),width=1000,height=800)

  m <- matrix(c(1,2,3,4,5,5),nrow = 6,ncol = 1,byrow = TRUE)

  par(oma=c(1,1,1,1))

  layout(mat = m,heights = c(0.050,0.3,0.3,0.3,0.025,0.025))

  # title
  par(mar = c(0,0,0,4))
  plot(1, type = "n", axes=FALSE, xlab="", ylab="")
  ##title(main=paste("Fitted/Residuals ",title.model,title.fitted,sep=" - "),
  ## col.main='red', font.main=4, cex=1.5, outer=FALSE)

  # fitted vs true values
  par(mar = c(4,4,1,0))
  plot(true.value,fitted.value,
  col=as.numeric(display.data),bg=as.numeric(display.data),
  type='p', pch=21,
  xlab=title.fitted, ylab=paste("Fitted - ",title.fitted," - ",title.model,sep="") )

```

```

title(main=paste("Fitted Vs True Value",title.fitted,sep=""),
col.main='blue', font.main=4,outer=FALSE)

# a scatter plot of Mean(Y) vs Standardized Residuals
par(mar = c(4,4,1,0))
plot(data$muhat,std.residuals,
col=as.numeric(display.data),bg=as.numeric(display.data),
type='p', pch=21,
xlab="Mean(Y)",
ylab=paste("Standardized Residuals for ",title.fitted,sep=""))

title(main=paste("Mean(Y) Vs Standardized Residuals by ",title.legend,sep=""),
col.main='blue', font.main=4,outer=FALSE)

# a scatter plot of Variance(Y) vs Standardized Residuals
par(mar = c(4,4,1,0))
plot(data$Vhat,std.residuals,
col=as.numeric(display.data),bg=as.numeric(display.data),
type='p', pch=21,
xlab="Variance(Y)",
ylab=paste("Standardized Residuals for ",title.fitted,sep=""))

title(main=paste("Variance(Y) Vs Standardized Residuals by ",title.legend,sep=""),
col.main='blue', font.main=4,outer=FALSE)

# legend
par(mar = c(0,0,0,0))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
legend(x="bottom", inset=0, paste(title.legend,levels(as.factor(muqt))),
fill=levels(as.factor(muqt)) ,cex=1, horiz = TRUE)

dev.off()

}

```

```

#-----
# Read the data
#-----

## Read Training Data
midtermtrain <- read.table(file = "http://www2.isye.gatech.edu/~ymei/7406/midtermtrain.c

## Testing Data
midtermtest  <- read.table(file = "http://www2.isye.gatech.edu/~ymei/7406/midtermtest.cs

#-----
# Data preparation
#-----
## Some plots for exploratory data analysis
X1 <- midtermtrain[,1]
X2 <- midtermtrain[,2]
muhat <- apply(midtermtrain[,3:202], 1, mean)
Vhat  <- apply(midtermtrain[,3:202], 1, var)

## regression with poly terms
poly.x1.max <- 6
poly.x2.max <- 6

X1_poly <- poly(X1,poly.x1.max,raw=TRUE)[-1]
X2_poly <- poly(X2,poly.x2.max,raw=TRUE)[-1]

```



```

data0 <- data.frame(X1 = X1,
X2=X2,
X1_poly,
X2_poly,
muhat = muhat,
Vhat = Vhat)

# muhat and vhat quartiles
muqt <- as.integer(cut(muhat, quantile(muhat, probs=0:4/4), include.lowest=TRUE))
Vqt <- as.integer(cut(Vhat, quantile(Vhat, probs=0:4/4), include.lowest=TRUE))

#-----
# Data Exploration and Analysis
#-----
## dim=2911x202
## The first two columns are X1 and X2 values, and the last 200 columns are the Y values
dim(midtermtrain)

## This should be a 1066*2 matrix
## Please add two columns for your estimation of the mean and variance of the Y variable.
dim(midtermtest)

#-----
# Box and Density Plots - Data Exploration and Analysis
#-----

png("mt_train_data_anal_1.png",width=1000,height=800)

m <- matrix(c(1,2,3,4),nrow = 4,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(0.05,0.35,0.3,0.3))

par(mar = c(0,0,0,5))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
##title(main='Training Data Analysis', col.main='red', font.main=4, cex=1.5, outer=TRUE)

```

```

par(mar = c(4,4,1,0))

boxplot(cbind(X1=midtermtrain[,1], X2=midtermtrain[,2], Y=stack(midtermtrain[,3:202])[,1],
"Mean(Y)"=muhat, "Variance(Y)"=Vhat),
col=c("gold", "darkgreen"),
main="Boxplot X1,X2,Y, Mean(Y) and Variance(Y) ",
col.main='red', font.main=4)

par(mar = c(4,4,1,0))

plot(density(muhat), main="Density Plot - Main(Y)")
polygon(density(muhat), col="red", border="blue")

par(mar = c(4,4,1,0))

plot(density(Vhat), main="Density Plot - Variance(Y)")
polygon(density(Vhat), col="red", border="blue")

dev.off()


png("mt_train_data_anal_density_plot.png", width=1000, height=800)

m <- matrix(c(1,2,3,4), nrow = 4, ncol = 1, byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m, heights = c(0.25,0.25,0.25,0.25))

for (i in c(100,1000,1500,2500)) {
  par(mar = c(4,4,1,0))

  y.density <- density(stack(midtermtrain[i,3:202])[,1])

  plot(y.density, main=paste("Density Plot of Y for X1=", midtermtrain[i,1], "X2=", midtermtrain[i,2], "Y=",

```

```

polygon(y.density,col="red", border="blue")
}

dev.off()

#-----
# Correlation Plots - Data Exploration and Analysis
#-----
# The correlation table (the last column is Y)
corr=round(cor( data0[,c("X1","X2","muhat","Vhat")] ),2)
corr
png("mt_corr_plot.png",width=1000,height=800)

m <- matrix(c(1,2),nrow = 2,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(0.5,0.5))

par(mar = c(4,4,1,0))

corrplot(corr, order = "AOE", cl.ratio = 0.2, cl.align = "r",
tl.pos = "d",tl.srt = 60)
title(main="Correlation Plot - X1,X2,Mean,Variance",
col.main='red', font.main=4,outer=FALSE)

dev.off()

#-----
# Matrix Scatter Plots
#-----

## scatter plot for possible collinrarity
png("mt_splom_scatter_matrix.png",width=1000,height=800)
splom( data0[,c("X1","X2","muhat","Vhat")] , pscales = 0,main="Matrix Scatter Plot",
col.main='red', font.main=4, xlab="")
dev.off()

```

```

par(mar = c(4,4,1,0))

png("mt_ggp_corr_plot.png",width=1000,height=800)

ggpairs(data0[,c("X1","X2","muhat","Vhat")],
title = "Correlation Plot",
upper = list (
mapping = ggplot2::aes(size = 16),
color = 'red'
),
lower = list(
continuous = "smooth",
combo = "facetdensity",
mapping = ggplot2::aes(color = muhat)
),
axisLabels='show')

dev.off()

#-----
# X1,X2 - muhat,Vhat Plots
#-----

png("mt_scatter_detailed_plot.png",width=1000,height=800)

par(mfrow = c(2,2))

## Or you can first create an initial plot of one line
##           and then iteratively add the lines
##
## below is an example to plot X1 vs. muhat for different X2 values
##
flag <- which(data0$X2 == 0)
plot(data0$X1[flag], data0$muhat[flag], type="l", xlim=range(data0$X1), ylim=range(data0$muhat),
xlab="X1", ylab="Mean(Y)",col="blue")
for (j in 1:40){
  flag <- which(data0$X2 == 0.1*j)
  lines(data0$X1[flag], data0$muhat[flag])
}

```

```

title(main="X1 Vs Mean(Y) - for different X2",
col.main='red', font.main=4,outer=FALSE)

## Or you can first create an initial plot of one line
##      and then iteratively add the lines
##
##  below is an example to plot X2 vs. muhat for different X1 values
##
flag <- which(data0$X1 == 0)
plot(data0$X2[flag], data0$muhat[flag], type="l", xlim=range(data0$X2), ylim=range(data0$muhat),
xlab="X2", ylab="Mean(Y)",col="blue")
for (j in 1:70){
  flag <- which(data0$X1 == 0.1*j)
  lines(data0$X2[flag], data0$muhat[flag])
}
title(main="X2 Vs Mean(Y) - for different X1",
col.main='red', font.main=4,outer=FALSE)

# variance vs X1 for different values of X2
flag <- which(data0$X2 == 0)
plot(data0$X1[flag], data0$Vhat[flag], type="l", xlim=range(data0$X1), ylim=range(data0$Vhat),
xlab="X1", ylab="Variance(Y)",col="red")
for (j in 1:40){
  flag <- which(data0$X2 == 0.1*j)
  lines(data0$X1[flag], data0$Vhat[flag])
}
title(main="X1 Vs Variance(Y)- for different X2",
col.main='red', font.main=4,outer=FALSE)

# variance vs X2 for different values of X1
flag <- which(data0$X1 == 0)
plot(data0$X2[flag], data0$Vhat[flag], type="l", xlim=range(data0$X2), ylim=range(data0$Vhat),
xlab="X2", ylab="Variance(Y)",col="red")
for (j in 1:70){
  flag <- which(data0$X1 == 0.1*j)
  lines(data0$X2[flag], data0$Vhat[flag])
}
title(main="X2 Vs Variance(Y) - for different X1",

```

```

col.main='red', font.main=4,outer=FALSE)

dev.off()

#-----
# X1 vs X2 scatter plots for mean and variance
#-----

png("mt_matplot_x1_x2_mu_v_qt.png",width=1000,height=800)

m <- matrix(c(1,2,3,4,5),nrow = 5,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(.05,0.375,.05,0.375,.05))

par(mar = c(0,0,0,0))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
##title(main="X1 X2 Vs Mean(Y)/Variance(Y) Quartile Bands", col.main='red', font.main=4

par(mar = c(4,4,1,0))
plot(data0$X1,data0$X2,type='p', pch=21, col=as.numeric(muqt),
bg=as.numeric(muqt),
xlab="X1", ylab="X2")
title(main="Mean(Y) Quartiles", col.main='red', font.main=4 ,outer=FALSE)

par(mar = c(0,0,0,0))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")

par(mar = c(4,4,1,0))
plot(data0$X1,data0$X2,type='p', pch=21, col=as.numeric(Vqt),
bg=as.numeric(Vqt),
xlab="X1", ylab="X2")
title(main="Variance(Y) Quartiles", col.main='red', font.main=4,outer=FALSE)

par(mar = c(0,0,0,0))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
legend(x="bottom", inset=0, paste("Quartile",levels(as.factor(Vqt))),
fill=levels(as.factor(muqt)) ,cex=1, horiz = TRUE)

```

```
dev.off()
```

```
#-----  
# Muhat to variance plot, is the variance uniform ( required for ols method)  
# If variance is different heterodascacity - need to look at logit which  
# does not assume homodascacity  
# Not sure if nomality is met here too  
#-----  
png("mt_matplot_mu_v_qt.png",width=1000,height=800)  
  
m <- matrix(c(1,2,3),nrow = 3,ncol = 1,byrow = TRUE)  
  
par(oma=c(1,1,1,1))plot  
  
layout(mat = m,heights = c(.05,0.9,.05))  
  
par(mar = c(0,0,0,0))  
plot(1, type = "n", axes=FALSE, xlab="", ylab="")  
##title(main="Mean(Y) Vs Variance(Y)", col.main='red', font.main=4 ,outer=TRUE)  
  
par(mar = c(4,4,1,0))  
plot(data0$muhat,data0$Vhat,type='p', pch=21, col=as.numeric(Vqt),  
bg=as.numeric(Vqt),  
xlab="Mean(Y) - muhat", ylab="Variance(Y) - Vhat")  
title(main="Mean(Y) Vs Variance(Y) ", col.main='red', font.main=4 ,outer=FALSE)  
  
par(mar = c(0,0,0,0))  
plot(1, type = "n", axes=FALSE, xlab="", ylab="")  
legend(x="bottom", inset=0, paste("Variance(Y) Quartile",levels(as.factor(Vqt))),  
fill=levels(as.factor(muqt)) ,cex=1, horiz = TRUE)  
  
dev.off()
```

```

#-----
# Fitting Models - Training and Cross Validation
#-----

# Keep track of MSE errors for different models
muhat.models <- c ('Logistic Regression','Logistic Regression(Poly)',
'Linear Regression','Linear Regression(Poly)','LASSO')

vhat.models <- c ('Logistic Regression','Logistic Regression(Poly)',
'Linear Regression','Linear Regression(Poly)','LASSO')

error.type <- c('TRAIN_MSE','TEST_MSE')

mse.models <- matrix(NA,length(muhat.models),length(error.type))
rownames(mse.models) <- muhat.models
colnames(mse.models) <- error.type

v.models <- matrix(NA,length(vhat.models),length(error.type))
rownames(v.models) <- vhat.models
colnames(v.models) <- error.type

# Split the training data into training and test
validation.test.percent = 5
test.count <- round(dim(data0)[1] * (validation.test.percent/100))
test.rows <- sort(sample(1:dim(data0)[1],test.count,replace=FALSE))

data0.test <- data0[test.rows,]
data0.train <- data0[-test.rows,]

muqt.test <- muqt[test.rows]
Vqt.test <- Vqt[test.rows]
muqt.train <- muqt[-test.rows]
Vqt.train <- Vqt[-test.rows]

# 10 Fold Cross Validation - Folds
#
nfolds <- 10
folds <- sample(1:nfolds,length(data0.train$X1),replace=TRUE)

```



```

#-----
# Fitting muhat
#
#-----

#-----
# muhat - Standard Logistic regression
#-----
# train using all train data

# train using all train data for minimum poly
lr.mu.fit <- glm(muhat~X1*X2,
data=data0.train,
family=binomial(logit))

train.residuals <- data0.train$muhat-lr.mu.fit$fitted
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lr.mu.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$muhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

mse.models["Logistic Regression",1] <- mean(train.residuals^2)
mse.models["Logistic Regression",2] <- mean(test.residuals^2)

summary(lr.mu.fit)

# chi sq test
chi.lr.muhat <- chisq.test(data0.train$muhat, lr.mu.fit$fitted)
print(chi.lr.muhat)

#-----
# muhat - Plot residuals vs X1,X2 by mean
#-----
plotxy.residuals("mt_rse_plot_mean_trg_lr",
data0.train,
lr.mu.fit$residuals,

```

```

train.stdres,
muqt.train,
data0.train$muhat,
lr.mu.fit$fitted,
"Mean(Y)",
"Mean(Y) Quartile",
"Logistic Regression"
)

```

```

#-----
# muhat - Poly Logistic regression
#-----
# train using all train data
# use 10 fold cross validation to choose the best poly term
poly.mse <- matrix(NA,poly.x1.max,poly.x2.max)
rownames(poly.mse) <- paste('X1_',c(1:6),sep="")
colnames(poly.mse) <- paste('X2_',c(1:6),sep="")

for ( p.x1 in 1:poly.x1.max ) {

for ( p.x2 in 1:poly.x2.max ) {

pred.mse <- matrix(NA,nfolds,1)

for ( i in 1:nfolds) {

lrp.mu.fit <- glm(muhat~poly(X1,p.x1)*poly(X2,p.x2),data=data0.train[folds!=i,],
family=binomial(logit))

lrp.pred <- predict(lrp.mu.fit,data0.train[folds==i,c("X1","X2")],type='response')

pred.mse[i,1] <- mean((lrp.pred-data0.train$muhat[folds==i])^2)
}

poly.mse[p.x1,p.x2] <- apply(pred.mse,2,mean)
}
}

```

```

# minimum poly with complexity factored in by multiplying with log of poly+1
poly.lrp.min <- arrayInd(which.min(poly.mse),dim(poly.mse))

poly.lrp.muhat.min.x1 <- poly.lrp.min[1]
poly.lrp.muhat.min.x2 <- poly.lrp.min[2]

# train using all train data for minimum poly
lrp.mu.fit <- glm(muhat~poly(X1,poly.lrp.muhat.min.x1)*poly(X2,poly.lrp.muhat.min.x2),
data=data0.train,
family=binomial(logit))

train.residuals <- data0.train$muhat-lrp.mu.fit$fitted
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lrp.mu.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$muhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

mse.models["Logistic Regression(Poly)",1] <- mean(train.residuals^2)
mse.models["Logistic Regression(Poly)",2] <- mean(test.residuals^2)

summary(lrp.mu.fit)

# chi sq test
chi.lrp.muhat <- chisq.test(data0.train$muhat, lrp.mu.fit$fitted)
print(chi.lrp.muhat)

#-----
# muhat - Plot residuals vs X1,X2 by mean
#-----
plotxy.residuals("mt_rse_plot_mean_trg_lr_poly",
data0.train,
lrp.mu.fit$residuals,
train.stdres,
muqt.train,
data0.train$muhat,
lrp.mu.fit$fitted,
"Mean(Y)",
"Mean(Y) Quartile",

```

```

"Logistic Regression(Poly)"
)

#-----
# Muhat to variance plot, is the variance uniform ( required for ols method)
# If variance is different heterodescacity - need to look at logit which
# does not assume homodescacity
# Not sure if nomality is met here too
#-----

rocol.mse <- arrayInd(c(1:length(c(poly.mse))),dim(poly.mse))
cv.labels <- paste(rownames(poly.mse)[rocol.mse[,1]],colnames(poly.mse)[rocol.mse[,2]])

png("mt_cvplot_mu_lg_poly.png",width=1000,height=800)

m <- matrix(c(1,2,3,3),nrow = 4,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(.05,0.80,.05,0.1))

par(mar = c(0,0,0,0))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
##title(main="Logistic Regression Mean(Y)- Cross Validation MSE plot Vs Poly Terms",
## col.main='red', font.main=4 ,outer=TRUE)

par(mar = c(4,4,1,0))

plot(c(1:length(c(poly.mse))),c(poly.mse),xaxt='n',type='p',pch=21,
xlab="Poly Terms", ylab="Cross Validation MSE - muhat", bg='blue')

lo <- loess(c(poly.mse)~c(1:length(c(poly.mse))))

lines(predict(lo), col='red', lwd=2)

axis(1,at=1:36,labels=cv.labels,cex=.5)

par(mar = c(1,1,1,1))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")

```

```

legend(x="bottom", inset=0,
c("MSE from Logistic Regression Cross for Mean(Y)","Smoothing Line for MSE"),
fill=c("blue","red") ,
cex=1, horiz = FALSE)

```

```

dev.off()

```

```

#-----
# muhat - Standard linear regression
#-----
# train using all train data
lm.mu.fit <- lm(muhat~X1+X2,data=data0.train)

train.residuals <- data0.train$muhat-lm.mu.fit$fitted.values
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lm.mu.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$muhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

mse.models['Linear Regression',1] <- mean(train.residuals^2)
mse.models['Linear Regression',2] <- mean(test.residuals^2)

summary(lm.mu.fit)

# chi sq test

```

```

chi.lm.muhat <- chisq.test(data0.train$muhat, lm.mu.fit$fitted.values)
print(chi.lm.muhat)

#-----
# muhat - Plot residuals vs X1,X2 by mean
#-----
plotxy.residuals("mt_rse_plot_mean_trg_slg",
data0.train,
lm.mu.fit$residuals,
train.stdres,
muqt.train,
data0.train$muhat,
lm.mu.fit$fitted.values,
"Mean(Y)",
"Mean(Y) Quartile",
"Linear Regression"
)

#-----
# muhat - linear regression - Poly
#-----
# use 10 fold cross validation to choose the best poly term
poly.mse <- matrix(NA,poly.x1.max,poly.x2.max)

for ( p.x1 in 1:poly.x1.max ) {

for ( p.x2 in 1:poly.x2.max ) {

pred.mse <- matrix(NA,nfolds,1)

for ( i in 1:nfolds) {

lmp.mu.fit <- lm(muhat~poly(X1,p.x1)+poly(X2,p.x2),data=data0.train[folds!=i,])

lmp.pred <- predict(lmp.mu.fit,data0.train[folds==i,c("X1","X2")])

pred.mse[i,1] <- mean((lmp.pred-data0.train$muhat[folds==i])^2)
}

poly.mse[p.x1,p.x2] <- apply(pred.mse,2,mean)

```

```

}
}

# minimum poly with complexity factored in by multiplying with log of poly+1
poly.min <- arrayInd(which.min(poly.mse),dim(poly.mse))

poly.muhat.min.x1 <- poly.min[1]
poly.muhat.min.x2 <- poly.min[2]

# train using all train data for minimum poly
lmp.mu.fit <- lm(muhat~poly(X1,poly.muhat.min.x1)+poly(X2,poly.muhat.min.x2),data=data0.train)

train.residuals <- data0.train$muhat-lmp.mu.fit$fitted.values
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lmp.mu.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$muhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

mse.models['Linear Regression(Poly)',1] <- mean(train.residuals^2)
mse.models['Linear Regression(Poly)',2] <- mean(test.residuals^2)

summary(lmp.mu.fit)

# chi squared test for test data
chi.poly.muhat <- chisq.test(data0.train$muhat, lmp.mu.fit$fitted.values)
print(chi.poly.muhat)

#-----
# muhat - Plot residuals vs X1,X2 by mean
#-----
plotxy.residuals("mt_rse_plot_mean_trg_lg_poly",
data0.train,
lmp.mu.fit$residuals,
train.stdres,
muqt.train,
data0.train$muhat,
lmp.mu.fit$fitted.values,
"Mean(Y)",

```

```

"Mean(Y) Quartile",
"Linear Regression Polynomial"
)

#-----
# muhat - Lasso
#-----
# cross validation to choose lambda for lasso on the bets poly model
#
#

lasso.lambda.grid <- 10^ seq (10,-6, length =1000)

poly.mse <- matrix(NA,poly.x1.max,poly.x2.max)

for ( p.x1 in 1:poly.x1.max ) {

  lasso.muhat.cols <- c(1:(2+p.x1-1))

  for ( p.x2 in 1:poly.x2.max ) {

    if ( p.x2 > 1 ) {
      lasso.muhat.cols <- c( lasso.muhat.cols,(2+poly.x1.max):(2+poly.x1.max+
p.x2-2))
    }

    pred.mse <- matrix(NA,nfolds,1)

    for ( i in 1:nfolds) {

      lasso.mu.fit <- cv.glmnet(as.matrix(data0.train[folds!=i,lasso.muhat.cols]),
data0.train$muhat[folds!=i],
alpha=1, lambda=lasso.lambda.grid)

      lasso.muhat.bestlam <- lasso.mu.fit$lambda.min

      lasso.pred=predict(lasso.mu.fit, s=lasso.muhat.bestlam ,
newx=as.matrix(data0.train[folds==i,lasso.muhat.cols]))

      pred.mse[i,1] <- mean((lasso.pred-data0.train$muhat[folds==i])^2)

```



```

}

poly.mse[p.x1,p.x2] <- apply(pred.mse,2,mean)
}
}

# minimum poly with complexity factored in by multiplying with log of poly+1
poly.min <- arrayInd(which.min(poly.mse),dim(poly.mse))

poly.lasso.muhat.min.x1 <- poly.min[1]
poly.lasso.muhat.min.x2 <- poly.min[2]


# CV the best poly term with lasso on the whole data to get the best lambda for it,
# using a wider range of lambda here
lasso.lambda.grid <- 10^ seq (10,-6, length =10000)
lasso.muhat.cols <- c(1:(2+poly.lasso.muhat.min.x1-1))
if ( poly.lasso.muhat.min.x2 > 1 ) {
lasso.muhat.cols <- c( lasso.muhat.cols,(2+poly.x1.max):(2+poly.x1.max+poly.lasso.muhat.min.x2
)}

# cross validate to get the best lambda
lasso.mu.fit <- cv.glmnet(as.matrix(data0.train[,lasso.muhat.cols]),data0.train$muhat,
alpha=1, lambda=lasso.lambda.grid)

lasso.muhat.bestlam <- lasso.mu.fit$lambda.min
coef(lasso.mu.fit, s = "lambda.min")

train.pred=predict(lasso.mu.fit, s=lasso.muhat.bestlam , newx=as.matrix(data0.train[,lasso.muhat.cols])
train.residuals <- data0.train$muhat-train.pred
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred=predict(lasso.mu.fit, s=lasso.muhat.bestlam , newx=as.matrix(data0.test[,lasso.muhat.cols])
test.residuals <- data0.test$muhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

mse.models['LASSO',1] <- mean(train.residuals^2)
mse.models['LASSO',2] <- mean(test.residuals^2)

```

```

summary(lasso.mu.fit)

tss <- sum((data0.train$muhat-mean(data0.train$muhat))^2)
rss <- sum(train.residuals^2)
lasso.muhat.R.squared <- (tss-rss)/tss
print(lasso.muhat.R.squared)

# chi squared test for test data
chi.lasso.muhat <- chisq.test(data0.train$muhat, train.pred)
print(chi.lasso.muhat)

#-----
# muhat - Plot residuals vs X1,X2 by mean
#-----
png("lasso_cv_plot_mean",width=1000,height=800)
plot(lasso.mu.fit)
title(main="Lasso CV Plot - Mean(Y)",
col.main='red', font.main=4,outer=FALSE)
dev.off()

plotxy.residuals("mt_rse_plot_mean_trg_lasso_poly",
data0.train,
lasso.mu.fit$residuals,
train.stdres,
muqt.train,
data0.train$muhat,
train.pred,
"Mean(Y)",
"Mean(Y) Quartile",
"Lasso Polynomial"
)

#-----
# Fit models for Vhat
#
#-----

```

```

#-----
# vhat - Standard Logistic regression
#-----
# train using all train data

# train using all train data for minimum poly
lr.v.fit <- glm(Vhat~X1*X2,
data=data0.train,
family=binomial(logit))

train.residuals <- data0.train$Vhat-lr.v.fit$fitted
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lr.v.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$muhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

v.models["Logistic Regression",1] <- mean(train.residuals^2)
v.models["Logistic Regression",2] <- mean(test.residuals^2)

summary(lr.v.fit)

# chi sq test
chi.lr.vhat <- chisq.test(data0.train$Vhat, lr.v.fit$fitted)
print(chi.lr.vhat)

#-----
# vhat - Plot residuals vs X1,X2 by mean
#-----
plotxy.residuals("mt_rse_plot_v_trg_lr",
data0.train,
lr.v.fit$residuals,
train.stdres,
Vqt.train,
data0.train$Vhat,
lr.v.fit$fitted,

```

```

"Variance(Y)",
"Variance(Y) Quartile",
"Logistic Regression"
)

#-----
# vhat - Poly Logistic regression
#-----
# train using all train data
# use 10 fold cross validation to choose the best poly term
poly.mse <- matrix(NA,poly.x1.max,poly.x2.max)
rownames(poly.mse) <- paste('X1_',c(1:6),sep="")
colnames(poly.mse) <- paste('X2_',c(1:6),sep="")

for ( p.x1 in 1:poly.x1.max ) {

for ( p.x2 in 1:poly.x2.max ) {

pred.mse <- matrix(NA,nfolds,1)

for ( i in 1:nfolds) {

lrp.v.fit <- glm(Vhat~poly(X1,p.x1)*poly(X2,p.x2),data=data0.train[folds!=i,],
family=binomial(logit))

lrp.pred <- predict(lrp.v.fit,data0.train[folds==i,c("X1","X2")],type='response')

pred.mse[i,1] <- mean((lrp.pred-data0.train$Vhat[folds==i])^2)
}

poly.mse[p.x1,p.x2] <- apply(pred.mse,2,mean)
}
}

# minimum poly with complexity factored in by multiplying with log of poly+1
poly.lrp.min <- arrayInd(which.min(poly.mse),dim(poly.mse))

poly.lrp.vhat.min.x1 <- poly.lrp.min[1]

```

```

poly.lrp.vhat.min.x2 <- poly.lrp.min[2]

# train using all train data for minimum poly
lrp.v.fit <- glm(Vhat~poly(X1,poly.lrp.vhat.min.x1)*poly(X2,poly.lrp.vhat.min.x2),
data=data0.train,
family=binomial(logit))

train.residuals <- data0.train$Vhat-lrp.v.fit$fitted
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lrp.v.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$Vhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

v.models["Logistic Regression(Poly)",1] <- mean(train.residuals^2)
v.models["Logistic Regression(Poly)",2] <- mean(test.residuals^2)

summary(lrp.v.fit)

# chi sq test
chi.lrp.vhat <- chisq.test(data0.train$Vhat, lrp.v.fit$fitted)
print(chi.lrp.vhat)

#-----
# vhat - Plot residuals vs X1,X2 by mean
#-----
plotxy.residuals("mt_rse_plot_v_trg_lr_poly",
data0.train,
lrp.v.fit$residuals,
train.stdres,
Vqt.train,
data0.train$Vhat,
lrp.v.fit$fitted,
"Variance(Y)",
"Variance(Y) Quartile",
"Logistic Regression(Poly)"
)

```

```

#-----
# Muhat to variance plot, is the variance uniform ( required for ols method)
# If variance is different heterodescacity - need to look at logit which
# does not assume homodescacity
# Not sure if nomality is met here too
#-----

rocol.mse <- arrayInd(c(1:length(c(poly.mse))),dim(poly.mse))
cv.labels <- paste(rownames(poly.mse)[rocol.mse[,1]],colnames(poly.mse)[rocol.mse[,2]])

png("mt_cvplot_v_lg_poly.png",width=1000,height=800)

m <- matrix(c(1,2,3,3),nrow = 4,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(.05,0.80,.05,0.1))

par(mar = c(0,0,0,0))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")

par(mar = c(4,4,1,0))

plot(c(1:length(c(poly.mse))),c(poly.mse),xaxt='n',type='p',pch=21,
xlab="Poly Terms", ylab="Cross Validation MSE - Vhat- Variance(Y)", bg='blue')

lo <- loess(c(poly.mse)~c(1:length(c(poly.mse))))

lines(predict(lo), col='red', lwd=2)

axis(1,at=1:36,labels=cv.labels,cex=.5)


par(mar = c(1,1,1,1))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
legend(x="bottom", inset=0,
c("MSE from Logistic Regression Cross for Variance(Y)","Smoothing Line for MSE"),
fill=c("blue","red") ,
cex=1, horiz = FALSE)

```

```
dev.off()
```

```
#-----  
# Vhat - Standard linear regression  
#-----
```

```

# train using all train data
lm.v.fit <- lm(Vhat~X1+X2,data=data0.train)

train.residuals <- data0.train$Vhat-lm.v.fit$fitted.values
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lm.v.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$Vhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

v.models['Linear Regression',1] <- mean(train.residuals^2)
v.models['Linear Regression',2] <- mean(test.residuals^2)

summary(lm.v.fit)

# chi squared test for test data
chi.lm.vhat <- chisq.test(data0.train$Vhat, lm.v.fit$fitted.values)
print(chi.lm.vhat)

#-----
# Plot residuals vs X1,X2 by Variance
#-----
plotxy.residuals("mt_rse_plot_var_v_trg_slg",
data0.train,
lm.v.fit$residuals,
train.stdres,
Vqt.train,
data0.train$Vhat,
lm.v.fit$fitted.values,
"Variance(Y)",
"Variance(Y) Quartile",
"Linear Regression"
)

#-----
# Vhat - linear regression - Poly

```



```

#-----
# use 10 fold cross validation to choose the best poly term
poly.mse <- matrix(NA,poly.x1.max,poly.x2.max)

for ( p.x1 in 1:poly.x1.max ) {

for ( p.x2 in 1:poly.x2.max ) {

pred.mse <- matrix(NA,nfolds,1)

for ( i in 1:nfolds) {

lmp.v.fit <- lm(Vhat~poly(X1,p.x1)+poly(X2,p.x2),data=data0.train[folds!=i,])

lmp.pred <- predict(lmp.v.fit,data0.train[folds==i,c("X1","X2")])

pred.mse[i,1] <- mean((lmp.pred-data0.train$Vhat[folds==i])^2)
}

poly.mse[p.x1,p.x2] <- apply(pred.mse,2,mean)
}
}

# minimum poly with complexity factored in by multiplying with log of poly+1
poly.min <- arrayInd(which.min(poly.mse),dim(poly.mse))

poly.Vhat.min.x1 <- poly.min[1]
poly.Vhat.min.x2 <- poly.min[2]

# train using all train data for minimum poly
lmp.v.fit <- lm(Vhat~poly(X1,poly.Vhat.min.x1)+poly(X2,poly.Vhat.min.x2),data=data0.train)

train.residuals <- data0.train$Vhat-lmp.v.fit$fitted.values
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred <- predict(lmp.v.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$Vhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

```

```

v.models['Linear Regression(Poly)',1] <- mean(train.residuals^2)
v.models['Linear Regression(Poly)',2] <- mean(test.residuals^2)

summary(lmp.v.fit)

# chi squared test for test data
chi.poly.vhat <- chisq.test(data0.train$Vhat, lmp.v.fit$fitted.values)
print(chi.poly.vhat)

#-----
# Vhat - Plot residuals vs X1,X2 by Variance
#-----
plotxy.residuals("mt_rse_plot_var_v_trg_lg_poly",
data0.train,
lmp.v.fit$residuals,
train.stdres,
Vqt.train,
data0.train$Vhat,
lmp.v.fit$fitted.values,
"Variance(Y)",
"Variance(Y) Quartile",
"Linear Regression Polynomial - Variance(Y)"
)

#-----
# Vhat - Lasso
#-----
# cross validation to choose lambda for lasso on the betspoly model
#
#

lasso.lambda.grid <- 10^ seq (10,-6, length =1000)

poly.mse <- matrix(NA,poly.x1.max,poly.x2.max)

for ( p.x1 in 1:poly.x1.max ) {

lasso.Vhat.cols <- c(1:(2+p.x1-1))

```

```

for ( p.x2 in 1:poly.x2.max ) {

  if ( p.x2 > 1 ) {
    lasso.Vhat.cols <- c( lasso.Vhat.cols,(2+poly.x1.max):(2+poly.x1.max+
p.x2-2))
  }

  pred.mse <- matrix(NA,nfolds,1)

  for ( i in 1:nfolds) {

    lasso.v.fit <- cv.glmnet(as.matrix(data0.train[folds!=i,lasso.Vhat.cols]),
data0.train$Vhat[folds!=i],
alpha=1, lambda=lasso.lambda.grid)

    lasso.Vhat.bestlam <- lasso.v.fit$lambda.min

    lasso.pred=predict(lasso.v.fit, s=lasso.Vhat.bestlam ,
newx=as.matrix(data0.train[folds==i,lasso.Vhat.cols]))

    pred.mse[i,1] <- mean((lasso.pred-data0.train$Vhat[folds==i])^2)
  }

  poly.mse[p.x1,p.x2] <- apply(pred.mse,2,mean)
}
}

# minimum poly with complexity factored in by muliplying with log of poly+1
poly.min <- arrayInd(which.min(poly.mse),dim(poly.mse))

poly.lasso.Vhat.min.x1 <- poly.min[1]
poly.lasso.Vhat.min.x2 <- poly.min[2]


# CV the best poly term with lasso on the whole data to get the best lambda for it,
# using a wider range of lambda here
lasso.lambda.grid <- 10^ seq (10,-6, length =10000)
lasso.Vhat.cols <- c(1:(2+poly.lasso.Vhat.min.x1-1))
if ( poly.lasso.Vhat.min.x2 > 1 ) {

```

```

lasso.Vhat.cols <- c( lasso.Vhat.cols,(2+poly.x1.max):(2+poly.x1.max+poly.lasso.Vhat.min
})

# cross validate to get the best lambda
lasso.v.fit <- cv.glmnet(as.matrix(data0.train[,lasso.Vhat.cols]),data0.train$Vhat,
alpha=1, lambda=lasso.lambda.grid)

lasso.Vhat.bestlam <- lasso.v.fit$lambda.min
coef(lasso.v.fit, s = "lambda.min")

train.pred=predict(lasso.v.fit, s=lasso.Vhat.bestlam , newx=as.matrix(data0.train[,lasso.V
train.residuals <- data0.train$Vhat-train.pred
train.sigma <- sqrt(sum(train.residuals^2)/(dim(data0.train)[1]-2-1))
train.stdres <- train.residuals / train.sigma

test.pred=predict(lasso.v.fit, s=lasso.Vhat.bestlam , newx=as.matrix(data0.test[,lasso.V
test.residuals <- data0.test$Vhat - test.pred
test.sigma <- sqrt(sum(test.residuals^2)/(dim(data0.test)[1]-2-1))
test.stdres <- test.residuals / test.sigma

v.models['LASSO',1] <- mean(train.residuals^2)
v.models['LASSO',2] <- mean(test.residuals^2)

summary(lasso.v.fit)

tss <- sum((data0.train$Vhat-mean(data0.train$Vhat))^2)
rss <- sum(train.residuals^2)
lasso.Vhat.R.squared <- (tss-rss)/tss
print(lasso.Vhat.R.squared)

# chi squared test for test data
chi.lasso.vhat <- chisq.test(data0.train$Vhat, train.pred)
print(chi.lasso.vhat)

#-----
# Vhat - Plot residuals vs X1,X2 by mean
#-----
png("lasso_cv_plot_v.png",width=1000,height=800)
plot(lasso.v.fit)
title(main="Lasso CV Plot - Variance(Y)",

```

```

col.main='red', font.main=3)
dev.off()

#-----
# Vhat Plot residuals vs X1,X2 by Variance
#-----
plotxy.residuals("mt_rse_plot_var_v_trg_lasso_poly",
data0.train,
lasso.v.fit$residuals,
train.stdres,
Vqt.train,
data0.train$Vhat,
train.pred,
"Variance(Y)",
"Variance(Y) Quartile",
"Lasso - Variance(Y)"
)

#-----
# Plot the MSE test error for different models
#-----
print(mse.models)

png("mt_train_test_mse_muhat_model_plot.png",width=1000,height=800)

matplot(mse.models, type="b", lty=1,xlab="Model", ylab="MSE(Mean(Y))",
col=seq_len(ncol(mse.models)),
cex.axis=1,xaxt="n")
axis(1, at = 1:length(muhat.models), labels = paste(muhat.models), cex.axis = .5)
title(main="Model Vs MSE(Mean(Y)) Training/Test",col.main='red', font.main=4)
legend("topright", colnames(mse.models),col=seq_len(ncol(mse.models)),cex=0.8,
fill=seq_len(ncol(mse.models)))

dev.off()

```

```

print(v.models)

png("mt_train_test_mse_v_model_plot.png",width=1000,height=800)

matplot(v.models, type="b", lty=1,xlab="Model", ylab="MSE(Variance(Y))",
col=seq_len(ncol(v.models)),
cex.axis=1,xaxt="n")
axis(1, at = 1:length(vhat.models), labels = paste(vhat.models), cex.axis = .5)
title(main="Model Vs MSE(Variance(Y)) Training/Test",col.main='red', font.main=4)
legend("topright", colnames(v.models),col=seq_len(ncol(v.models)),cex=0.8,
fill=seq_len(ncol(v.models)))

dev.off()

#-----
# Do a Bootstrap test to pick the best model statistically
#-----
### number of loops
B <- 100

### Final TE values for Mean(Y)
TEALL <- NULL

### Final TE values for Variance(y)
VEALL <- NULL

# create a matrix to hold the test MSE of the different models for each of the B cycles
TEALL <- matrix(NA,B,length(muhat.models),dimnames=list(1:B,muhat.models))
VEALL <- matrix(NA,B,length(vhat.models),dimnames=list(1:B,vhat.models));

# Split the training data into training and test
validation.test.percent <- 50

for (b in 1:B){

### randomly select 10% observations as testing data in each loop
test.count <- round(dim(data0)[1] * (validation.test.percent/100))
test.rows <- sort(sample(1:dim(data0)[1],test.count,replace=TRUE))

```

```

data0.test <- data0[test.rows,]
data0.train <- data0[-test.rows,]

muqt.test <- muqt[test.rows]
Vqt.test <- Vqt[test.rows]
muqt.train <- muqt[-test.rows]
Vqt.train <- Vqt[-test.rows]

#-----
# Logistic Regression
#-----
# muhat train using all train data
lr.mu.fit <- glm(muhat~X1*X2,
data=data0.train,
family=binomial(logit))

test.pred <- predict(lr.mu.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$muhat - test.pred
TEALL[b,"Logistic Regression"] <- mean(test.residuals^2)

# vhat train using all train data
lr.v.fit <- glm(Vhat~X1*X2,
data=data0.train,
family=binomial(logit))

test.pred <- predict(lr.v.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$Vhat - test.pred
VEALL[b,"Logistic Regression"] <- mean(test.residuals^2)

#-----
# Logistic Regression Poly
#-----
# muhat train using all train data
lrp.mu.fit <- glm(muhat~poly(X1,poly.lrp.muhat.min.x1)*poly(X2,poly.lrp.muhat.min.x2),
data=data0.train,
family=binomial(logit))

```

```

test.pred <- predict(lrp.mu.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$muhat - test.pred
TEALL[b,"Logistic Regression(Poly)"] <- mean(test.residuals^2)

# vhat train using all train data
lrp.v.fit <- glm(Vhat~poly(X1,poly.lrp.vhat.min.x1)*poly(X2,poly.lrp.vhat.min.x2),
data=data0.train,
family=binomial(logit))

test.pred <- predict(lrp.v.fit,data0.test[,c("X1","X2")],type='response')
test.residuals <- data0.test$Vhat - test.pred
VEALL[b,"Logistic Regression(Poly)"] <- mean(test.residuals^2)

#-----
# Linear Regression
#-----
# muhat train using all train data
lm.mu.fit <- lm(muhat~X1+X2,data=data0.train)
test.pred <- predict(lm.mu.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$muhat - test.pred
TEALL[b,"Linear Regression"] <- mean(test.residuals^2)

# Vhat train using all train data
lm.v.fit <- lm(Vhat~X1+X2,data=data0.train)
test.pred <- predict(lm.v.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$Vhat - test.pred
VEALL[b,"Linear Regression"] <- mean(test.residuals^2)

#-----
# Linear Regression Poly
#-----
# muhat train using all train data
lmp.mu.fit <- lm(muhat~poly(X1,poly.muhat.min.x1)+poly(X2,poly.muhat.min.x2),data=data0.
test.pred <- predict(lmp.mu.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$muhat - test.pred
TEALL[b,"Linear Regression(Poly)"] <- mean(test.residuals^2)

```



```

# Vhat train using all train data
lmp.v.fit <- lm(Vhat~poly(X1,poly.Vhat.min.x1)+poly(X2,poly.Vhat.min.x2),data=data0.train)
test.pred <- predict(lmp.v.fit,data0.test[,c("X1","X2")])
test.residuals <- data0.test$Vhat - test.pred
VEALL[b,"Linear Regression(Poly)"] <- mean(test.residuals^2)

#-----
#  Lasso
#-----
# muhat
lasso.mu.fit <- glmnet(as.matrix(data0.train[,lasso.muhat.cols]),data0.train$muhat,alpha=1,
lambda=lasso.muhat.bestlam)
test.pred <- predict(lasso.mu.fit, s=lasso.muhat.bestlam ,
newx=as.matrix(data0.test[,lasso.muhat.cols]))
test.residuals <- data0.test$muhat - test.pred
TEALL[b,"LASSO"] <- mean(test.residuals^2)

# Vhat
lasso.v.fit <- glmnet(as.matrix(data0.train[,lasso.Vhat.cols]),data0.train$Vhat,alpha=1,
lambda=lasso.Vhat.bestlam)
test.pred <- predict(lasso.v.fit, s=lasso.Vhat.bestlam ,
newx=as.matrix(data0.test[,lasso.Vhat.cols]))
test.residuals <- data0.test$Vhat - test.pred
VEALL[b,"LASSO"] <- mean(test.residuals^2)

}

#-----
#box plots of the Bootstrap B=100 run results
#-----

TM <- as.matrix(apply(TEALL, 2, mean))
TS <- as.matrix(apply(TEALL, 2, var))

VM <- as.matrix(apply(VEALL, 2, mean))
VS <- as.matrix(apply(VEALL, 2, var))

colnames(TM) <- "Average MSE(Mean(Y))"

```

```

colnames(TS) <- "Variance MSE(Mean(Y))"

colnames(VM) <- "Average MSE(Variance(Y))"
colnames(VS) <- "Variance MSE(Variance(Y))"

print(TM)
print(TS)

print(VM)
print(VS)

#-----
# PLOT
#-----
# muhat avergae MSE

png("mt_bootstrap_muhat_mse_var_plot.png",width=1000,height=800)

m <- matrix(c(1,2,3),nrow = 3,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(0.05,0.475,0.475))

par(mar = c(0,0,0,5))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
##title(main='BootStrap B=100 Results - Mean(Y)', col.main='red', font.main=4, cex=1.5,

par(mar = c(4,4,1,0))

matplot(TM, type="b", lty=1, xlab="Model", ylab="Average MSE(Mean(Y))",
col=seq_len(nrow(TM)),pch=1,xaxt="n")
axis(1, at = 1:length(muhat.models), labels = paste(muhat.models), cex.axis = 1)
title(main="Bootstrap Average MSE for Mean(Y)",
col.main='red', font.main=4)
legend("topright", colnames(TM),col=seq_len(ncol(TM)),cex=0.8,fill=seq_len(ncol(TM)))

## Variance
par(mar = c(4,4,1,0))

```

```

matplot(TS, type="b", lty=1, xlab="Model", ylab="Variance MSE(Mean(Y))",
col=seq_len(nrow(TS)),pch=1,xaxt="n")
axis(1, at = 1:length(muhat.models), labels = paste(muhat.models), cex.axis = 1)
title(main="Bootstrap Variance MSE for Mean(Y)",
col.main='red', font.main=4)
legend("topright", colnames(TS),col=seq_len(ncol(TS)),cex=0.8,fill=seq_len(ncol(TS)))

dev.off()

# Vhat avergae MSE

png("mt_bootstrap_vhat_mse_var_plot.png",width=1000,height=800)

m <- matrix(c(1,2,3),nrow = 3,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(0.05,0.475,0.475))

par(mar = c(0,0,0,5))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
##title(main='BootStrap B=100 Results - Variance(Y)', col.main='red', font.main=4, cex=1.5, ou

par(mar = c(4,4,1,0))

matplot(VM, type="b", lty=1, xlab="Model", ylab="Average MSE(Variance(Y))",
col=seq_len(nrow(VM)),pch=1,xaxt="n")
axis(1, at = 1:length(vhat.models), labels = paste(vhat.models), cex.axis = 1)
title(main="Bootstrap Average MSE for Variance(Y)",
col.main='red', font.main=4)
legend("topright", colnames(VM),col=seq_len(ncol(VM)),cex=0.8,fill=seq_len(ncol(VM)))

## Variance
par(mar = c(4,4,1,0))

matplot(VS, type="b", lty=1, xlab="Model", ylab="Variance MSE(Variance(Y))",
col=seq_len(nrow(VS)),pch=1,xaxt="n")

```

```

axis(1, at = 1:length(vhat.models), labels = paste(vhat.models), cex.axis = 1)
title(main="Bootstrap Variance MSE for Variance(Y)",
col.main='red', font.main=4)
legend("topright", colnames(VS),col=seq_len(ncol(VS)),cex=0.8,fill=seq_len(ncol(VS)))

dev.off()


##boxplot
png("mt_boxplot_bootstrap_muhat_results.png",width=1000,height=800)

m <- matrix(c(1,2,3,4),nrow = 4,ncol = 1,byrow = TRUE)

par(oma=c(1,1,1,1))

layout(mat = m,heights = c(0.05,0.45,0.05,0.45))

par(mar = c(0,0,0,5))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
##title(main='BootStrap B=100 Results', col.main='red', font.main=4, cex=1.5, outer=TRUE)

par(mar = c(4,4,1,0))
boxplot(TEALL,col=(c("gold","darkgreen")),main="Bootstrap MSE for Mean(Y) by Model",
xlab="Model",ylab="MSE Mean(Y)",
col.main='red', font.main=4)

par(mar = c(0,0,0,5))
plot(1, type = "n", axes=FALSE, xlab="", ylab="")

par(mar = c(4,4,1,0))
boxplot(VEALL,col=(c("gold","darkgreen")),main="Bootstrap MSE for Variance(Y) by Model",
xlab="Model",ylab="MSE Variance(Y)",
col.main='red', font.main=4)

dev.off()

```

```

#-----
# T test and W test to pick the best method
#-----
# create a matrix to hold the test MSE of the different models for each of the B cycles

## mhat
best.mhat.method <- rownames(TM)[which.min(TM)]
print(best.mhat.method)

t.test.table.mhat <- matrix(NA,2,length(muhat.models)-1,dimnames=list(c('T-Test','Wilcox Test',
muhat.models[muhat.models!=best.mhat.method])));

for ( m in muhat.models[muhat.models!=best.mhat.method] ) {

t.test.table.mhat[1,m] <- t.test(TEALL[,best.mhat.method],TEALL[,m],paired=TRUE,
conf.level=.95)$p.value
t.test.table.mhat[2,m] <- wilcox.test(TEALL[,best.mhat.method],TEALL[,m],paired=TRUE,
conf.level=.95)$p.value
}

print(t.test.table.mhat)


## Vhat
best.Vhat.method <- rownames(VM)[which.min(VM)]
print(best.Vhat.method)

t.test.table.vhat <- matrix(NA,2,length(vhat.models)-1,dimnames=list(c('T-Test','Wilcox Test',
vhat.models[vhat.models!=best.Vhat.method])));

for ( m in vhat.models[vhat.models!=best.Vhat.method] ) {

t.test.table.vhat[1,m] <- t.test(VEALL[,best.Vhat.method],VEALL[,m],paired=TRUE,
conf.level=.95)$p.value
t.test.table.vhat[2,m] <- wilcox.test(VEALL[,best.Vhat.method],VEALL[,m],paired=TRUE,
conf.level=.95)$p.value
}

print(t.test.table.vhat)

```

```

#-----
# Pick the best model and train it on the whole training data
#-----,

X1_poly <- poly(midtermtest[,1],poly.x1.max,raw=TRUE)[-1]
X2_poly <- poly(midtermtest[,2],poly.x2.max,raw=TRUE)[-1]

midtermtest.datapoly <- data.frame(X1=midtermtest[,1],
X2=midtermtest[,2],
X1_poly,
X2_poly)

# muhat train using all train data poly LM
#lmp.mu.fit <- lm(muhat~poly(X1,poly.muhat.min.x1)+poly(X2,poly.muhat.min.x2),data=data0)
#midtermtest.poly.mean <- predict(lmp.mu.fit,midtermtest.datapoly[,c("X1","X2")])

# Vhat train using all train data poly LM
#lmp.v.fit <- lm(Vhat~poly(X1,poly.Vhat.min.x1)+poly(X2,poly.Vhat.min.x2),data=data0)
#midtermtest.lmp.variance <- predict(lmp.v.fit,midtermtest.datapoly[,c("X1","X2")])

# T test and W test says Lasso is good too based on p-value
#lasso.v.fit <- glmnet(as.matrix(data0[,lasso.Vhat.cols]),data0$Vhat,alpha=1,
# lambda=lasso.Vhat.bestlam)

#midtermtest.lasso.variance <- predict(lasso.v.fit, s=lasso.Vhat.bestlam ,
# newx=as.matrix(midtermtest.datapoly[,lasso.Vhat.cols]))

# Logistic regression muhat
# muhat train using all train data
lrp.mu.fit <- glm(muhat~poly(X1,poly.lrp.muhat.min.x1)*poly(X2,poly.lrp.muhat.min.x1),
data=data0,
family=binomial(logit))

```

```

midtermtest.mean <- predict(lrp.mu.fit,midtermtest.datapoly[,c("X1","X2")],
type='response')

# vhat train using all train data
lrp.v.fit <- glm(Vhat~poly(X1,poly.lrp.vhat.min.x1)*poly(X2,poly.lrp.vhat.min.x1),
data=data0,
family=binomial(logit))

midtermtest.variance <- predict(lrp.v.fit,midtermtest.datapoly[,c("X1","X2")],
type='response')


#-----
# Write the results csv file
#-----
resultdata <- data.frame(midtermtest,
mean=format(round(midtermtest.mean,6),nsmall=6),
variance=format(round(midtermtest.variance,6),nsmall=6))

write.csv(resultdata, file = "midterm_results.csv",quote=FALSE)

```

## 9 References

- [1] Trevor Hastie, Robert Tibshirani, Jerome Friedman, *L<sup>A</sup>T<sub>E</sub>X: Elements of Statistical Learning*, Elements of Statistical Learning Ed. 2, 2009.